



ArgoCon

Lightning Talk: Configuring Volumes for Parallel Workflow Reads and Writes



Lukonde Mwila, Amazon Web Services



Tim Collins, Pipekit

Intro's

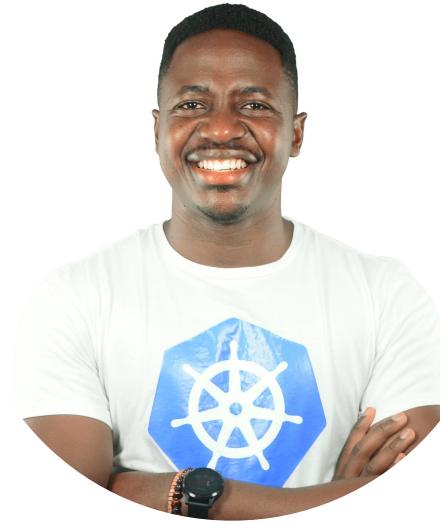


Tim Collins

Senior Infrastructure Engineer @ Pipekit.io



- Lives in the Argo Workflows Slack
 - Keen (but bad) kiteboarder



Lukonde Mwila

Senior Developer Advocate, Amazon EKS team @ AWS and a CNCF Ambassador



- Application development, solution architecture, cloud engineering, and DevOps workflows
- Cloud-native ecosystem contributor

The Problems

- ◆ How do I efficiently get from Workflows Step A to Workflow Step B?
- ◆ How do I efficiently share the same data with multiple Steps in parallel?
- ◆ How can I use semi-persistent data across many Workflows in parallel?

Persistent | Semi-Persistent | Transient

Storage Options

	Pros	Cons
S3	<ul style="list-style-type: none">• Easy setup• Artifacts visible in the UI	<ul style="list-style-type: none">• Files are tarred by Workflows before uploading - takes time and resources.• Copies are downloaded in the init container of each workflow task (so not RWX).• Only useful if you're in the cloud.
Minio	<ul style="list-style-type: none">• Easy Setup.• Artifacts visible in the UI• Useful on prem	<ul style="list-style-type: none">• Needs a disk behind it.• Small maintenance overhead.• Nuanced featureset differences with S3.• Files are tarred by Workflows before uploading - takes time and resources.• Copies are downloaded in the init container of each workflow task (so not RWX).
EBS	<ul style="list-style-type: none">• Easy setup• Some performance advantages over S3• No tarring required	<ul style="list-style-type: none">• Not RWX
EFS	<ul style="list-style-type: none">• RWX• No tarring required	<ul style="list-style-type: none">• Slow with parallel read/writes at scale• Only useful if you're in the cloud.

NFS-Server Provisioner

What is it?

<https://github.com/kubernetes-sigs/nfs-ganesha-server-and-external-provisioner>

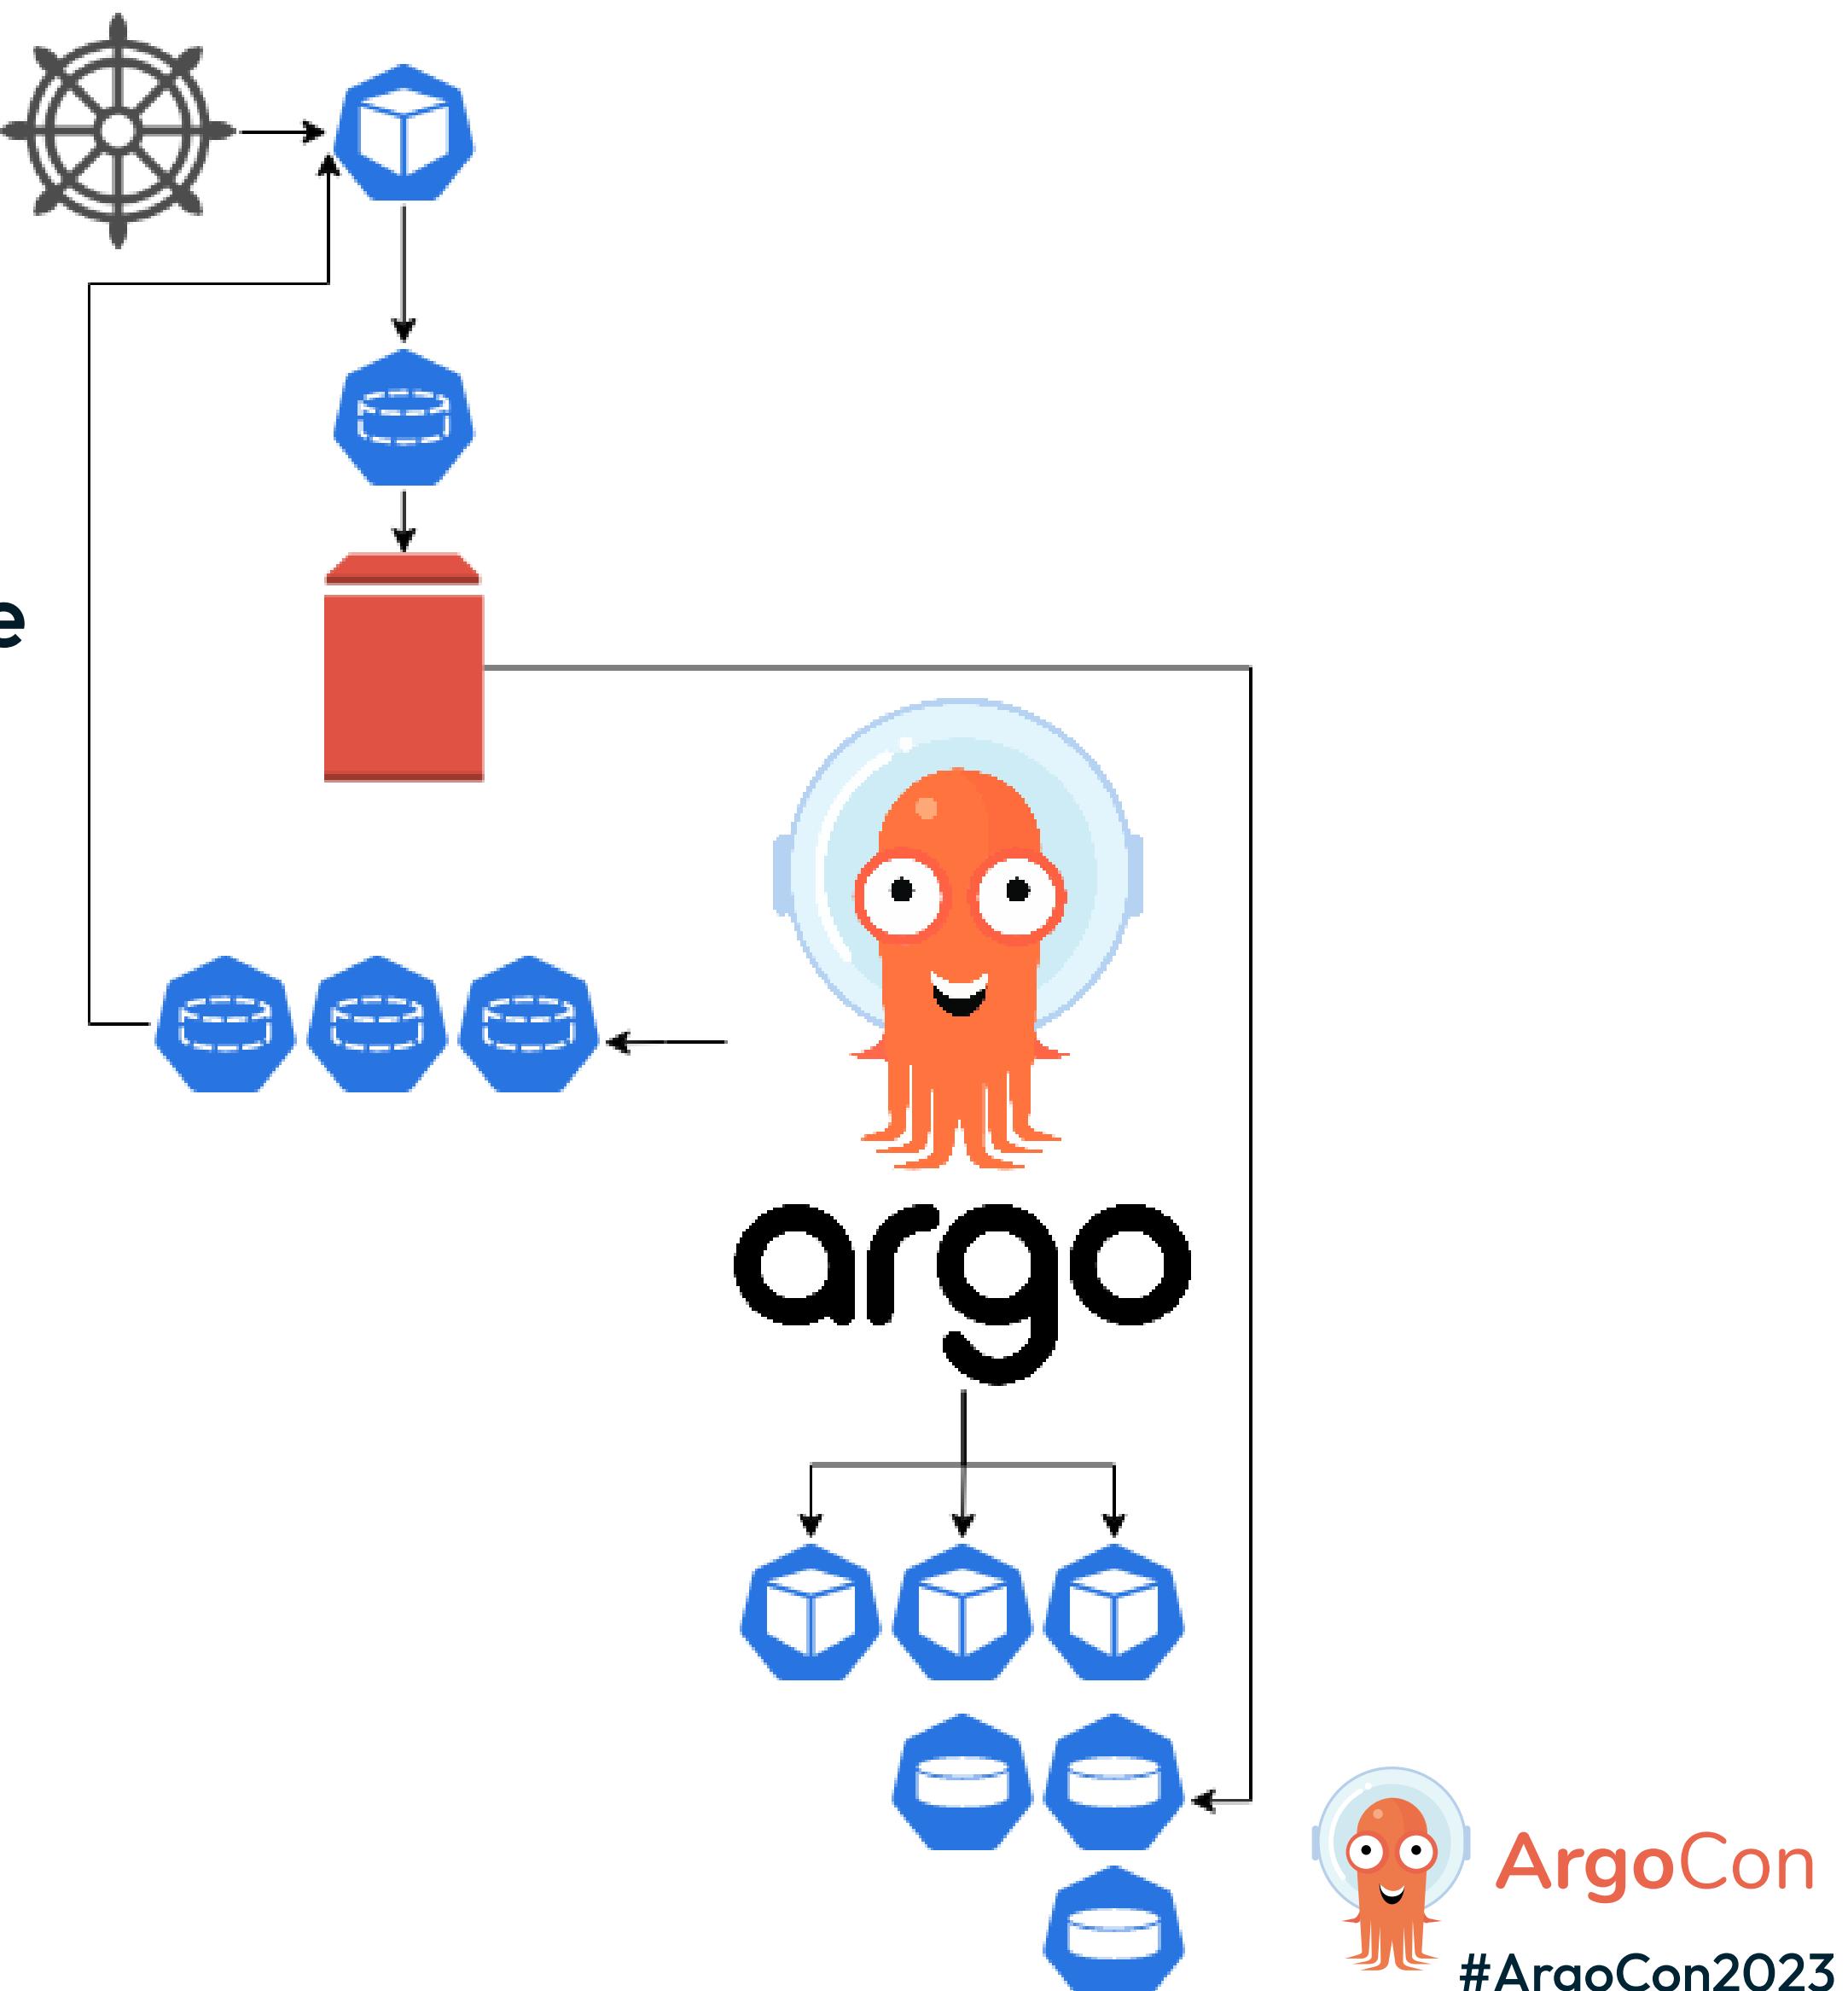
	Pros	Cons
NFS-server-provisioner	<ul style="list-style-type: none">• As EFS• Can be provisioned with a PVC• Useful on prem• Easy setup	<ul style="list-style-type: none">• Slow with parallel read/writes at scale.• Small maintenance overhead.



Full setup including ArgoCD deployment and a CI example using it can be found via
<https://github.com/pipekit/talk-demos>

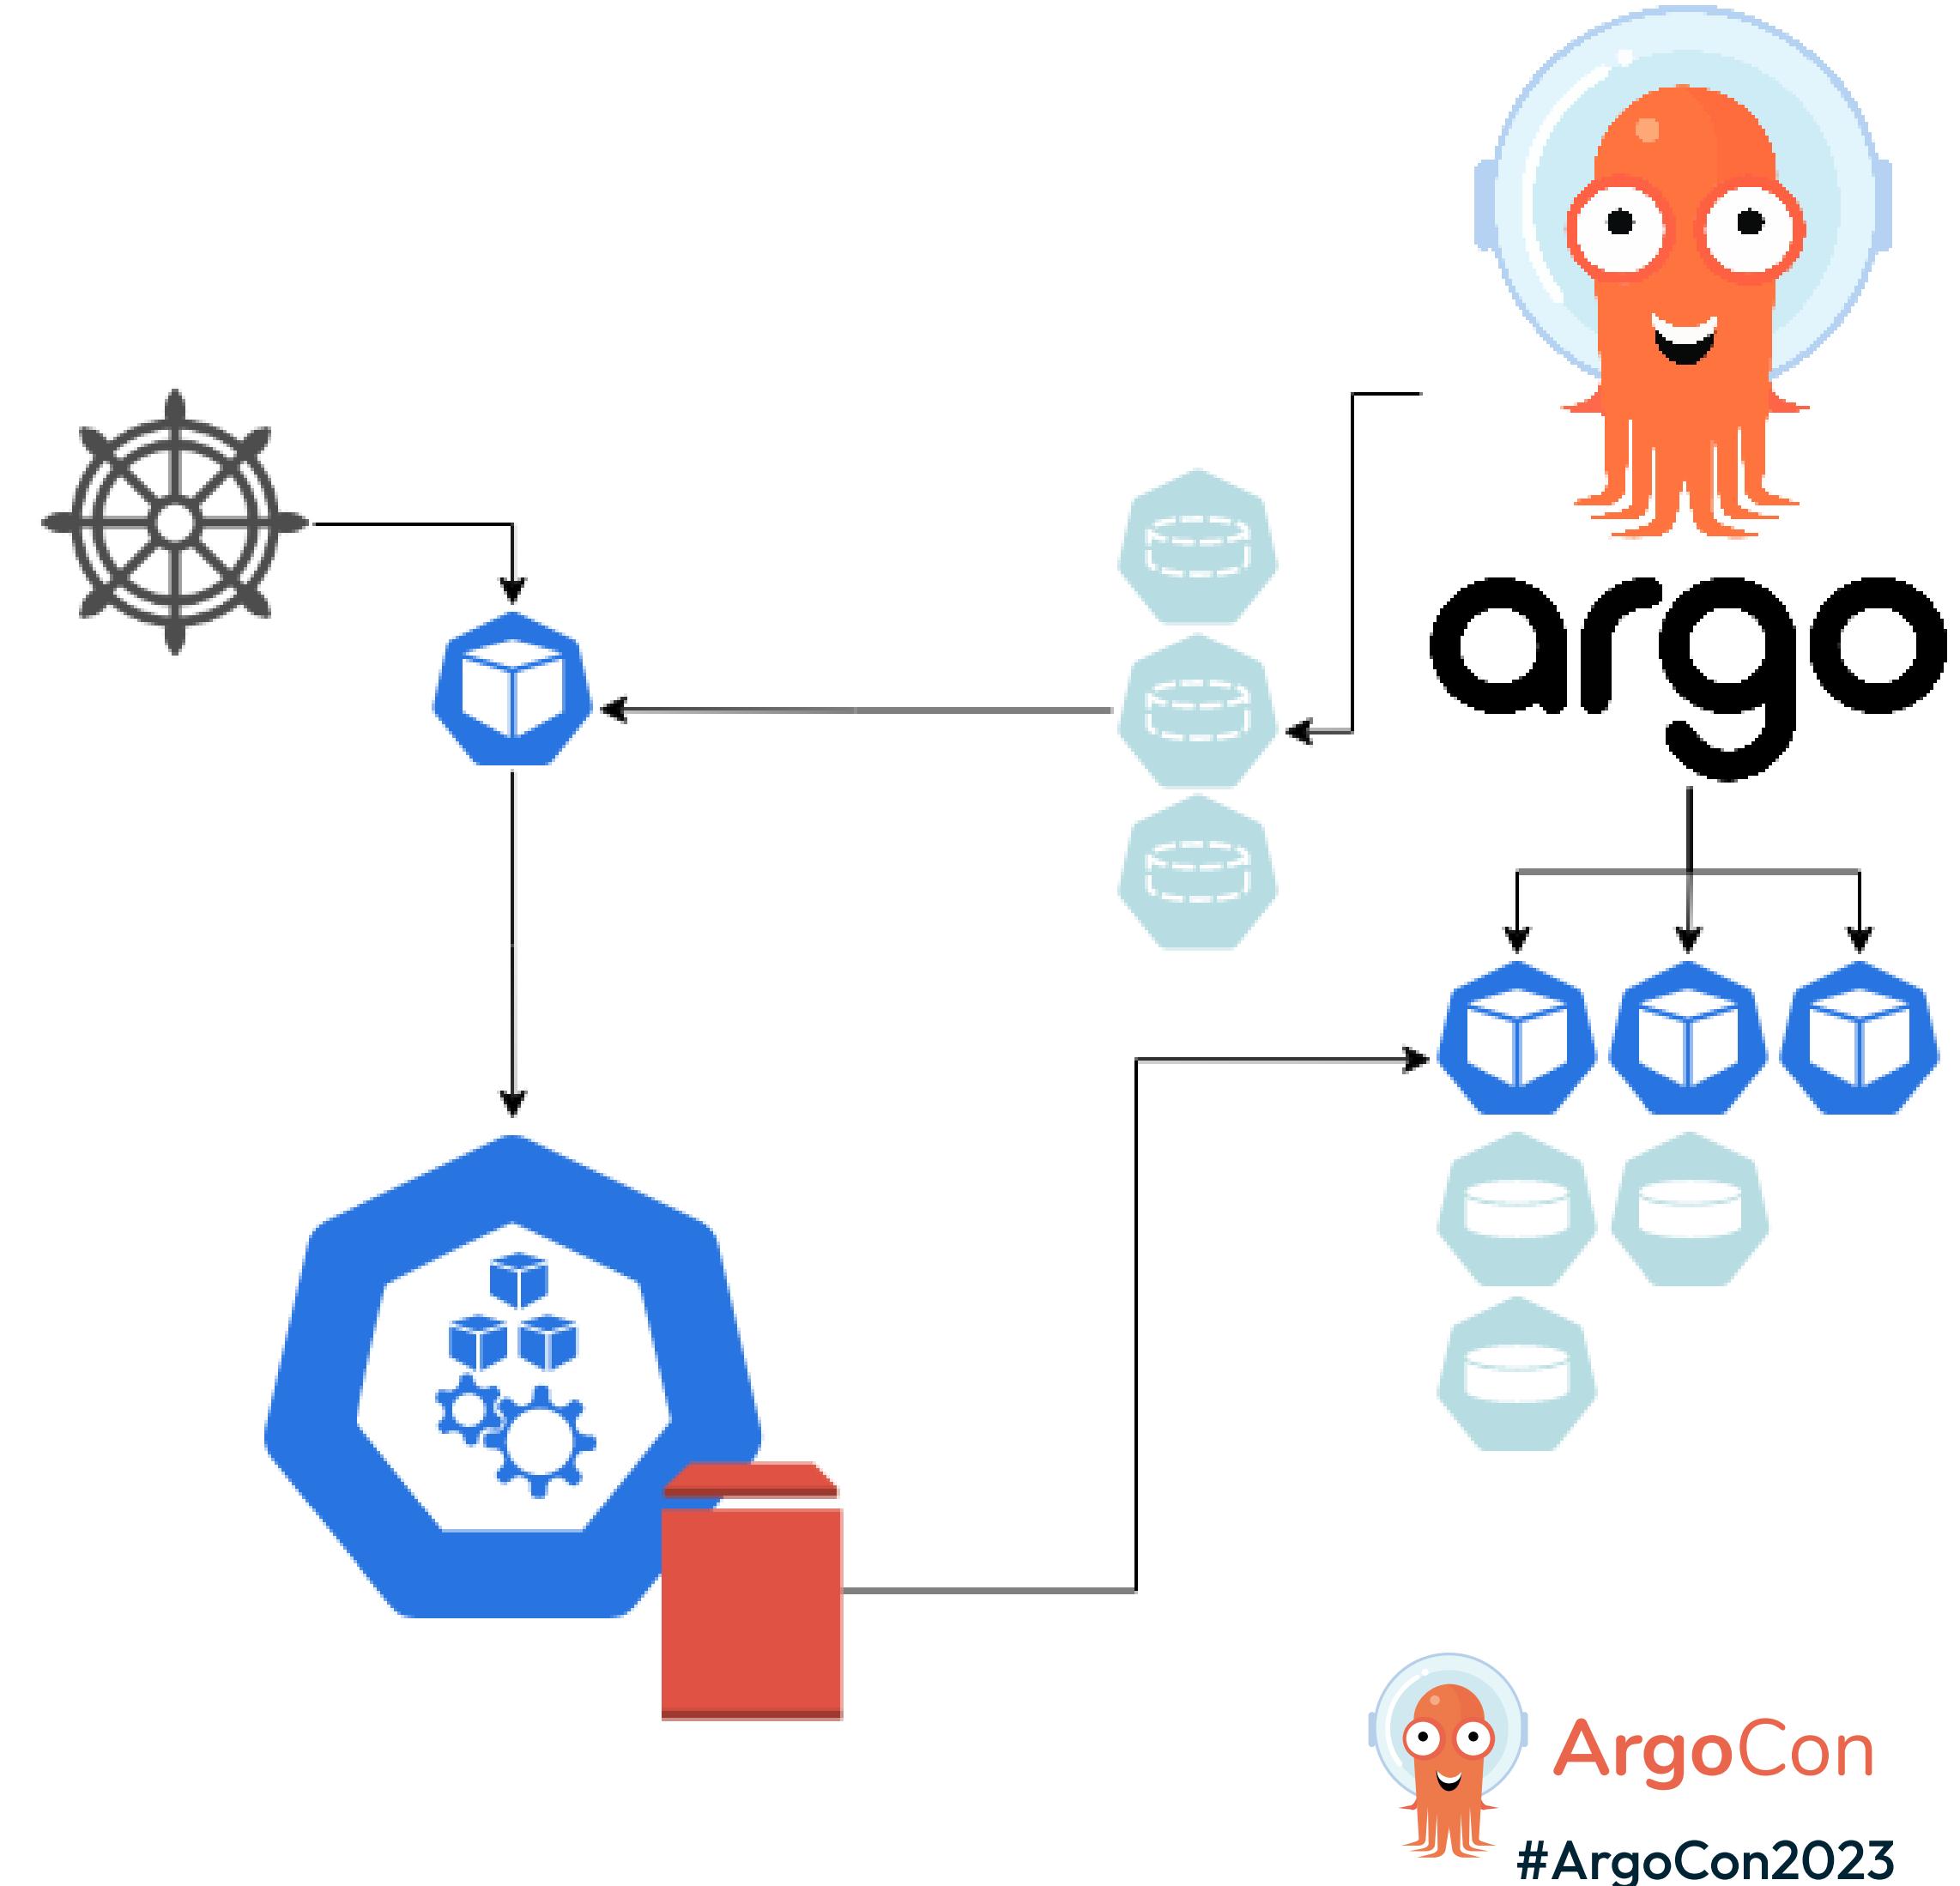
How does it work for semi-persistent data?

- ◆ Deployment creates an NFS Server Provisioner pod, PVC and storageClass.
- ◆ PVC provisions an EBS (RWO) disk and mounts it to the NFS Server Provisioner pod.
- ◆ We create our own PVCs using the new storageclass
- ◆ We are allocated Read-write-many persistent volumes, backed by EBS.
- ◆ In a Workflow, we mount these new volumes as required.



How does it work for transient data?

- Deployment creates an NFS Server Provisioner pod and storageClass.
- When we start a Workflow, we create volumeClaimTemplates to create PVCs that exist for the life of the Workflow.



```
1 apiVersion: argoproj.io/v1alpha1
2 kind: Workflow
3 metadata:
4- generateName: large-file-s3-
5 spec:
6   entrypoint: main
7
8   templates:
9     - name: main
10    dag:
11      tasks:
12        - name: example
13          template: example
14    container:
15      image: alpine
16      command:
17        - sh
18        - -c
19        - |
20-         cd /tmp
21           dd if=/dev/zero of=LARGEFILE bs=1 count=0 seek=10G
22   resources:
23     requests:
24       memory: 1Gi
25       cpu: 1
26   outputs:
27     artifacts:
28-       - name: large_file
29         path: /tmp/LARGEFILE
30
```

```
1 apiVersion: argoproj.io/v1alpha1
2 kind: Workflow
3 metadata:
4- generateName: large-file-nfs-
5 spec:
6   entrypoint: main
7+ volumeClaimTemplates:
8+
9+   - metadata:
10+     name: workdir
11+     spec:
12+       accessModes: [ "ReadWriteMany" ]
13+       storageClassName: nfs
14+       resources:
15+         requests:
16+           storage: 50Gi
16   templates:
17     - name: main
18     dag:
19       tasks:
20         - name: example
21           template: example
22     - name: example
23       container:
24         image: alpine
25         command:
26           - sh
27           - -c
28           - |
29+         cd /workdir
30           dd if=/dev/zero of=LARGEFILE bs=1 count=0 seek=10G
31   resources:
32     requests:
33       memory: 1Gi
34       cpu: 1
35+   volumeMounts:
36+     - name: workdir
37+       mountPath: /workdir
38
```

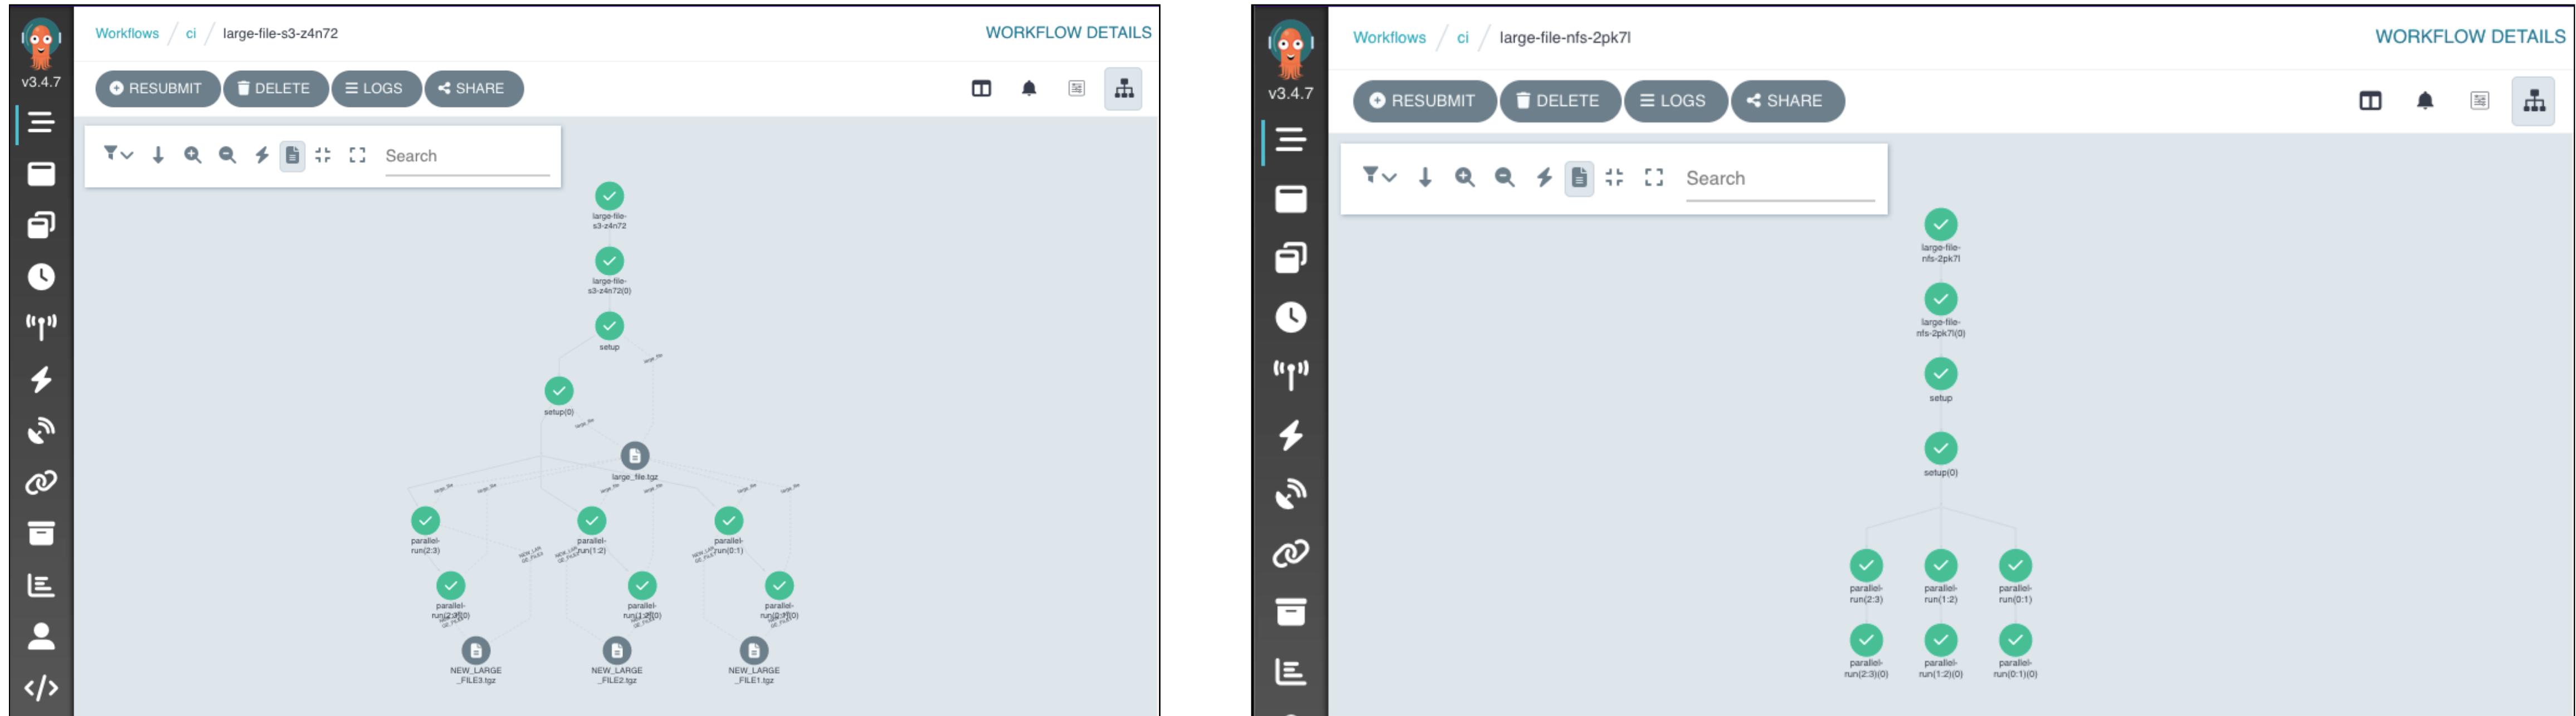
Comparing Cost

Storage Costs (50 GB/month)

Indicative figures for 50GB/month storage and transfer

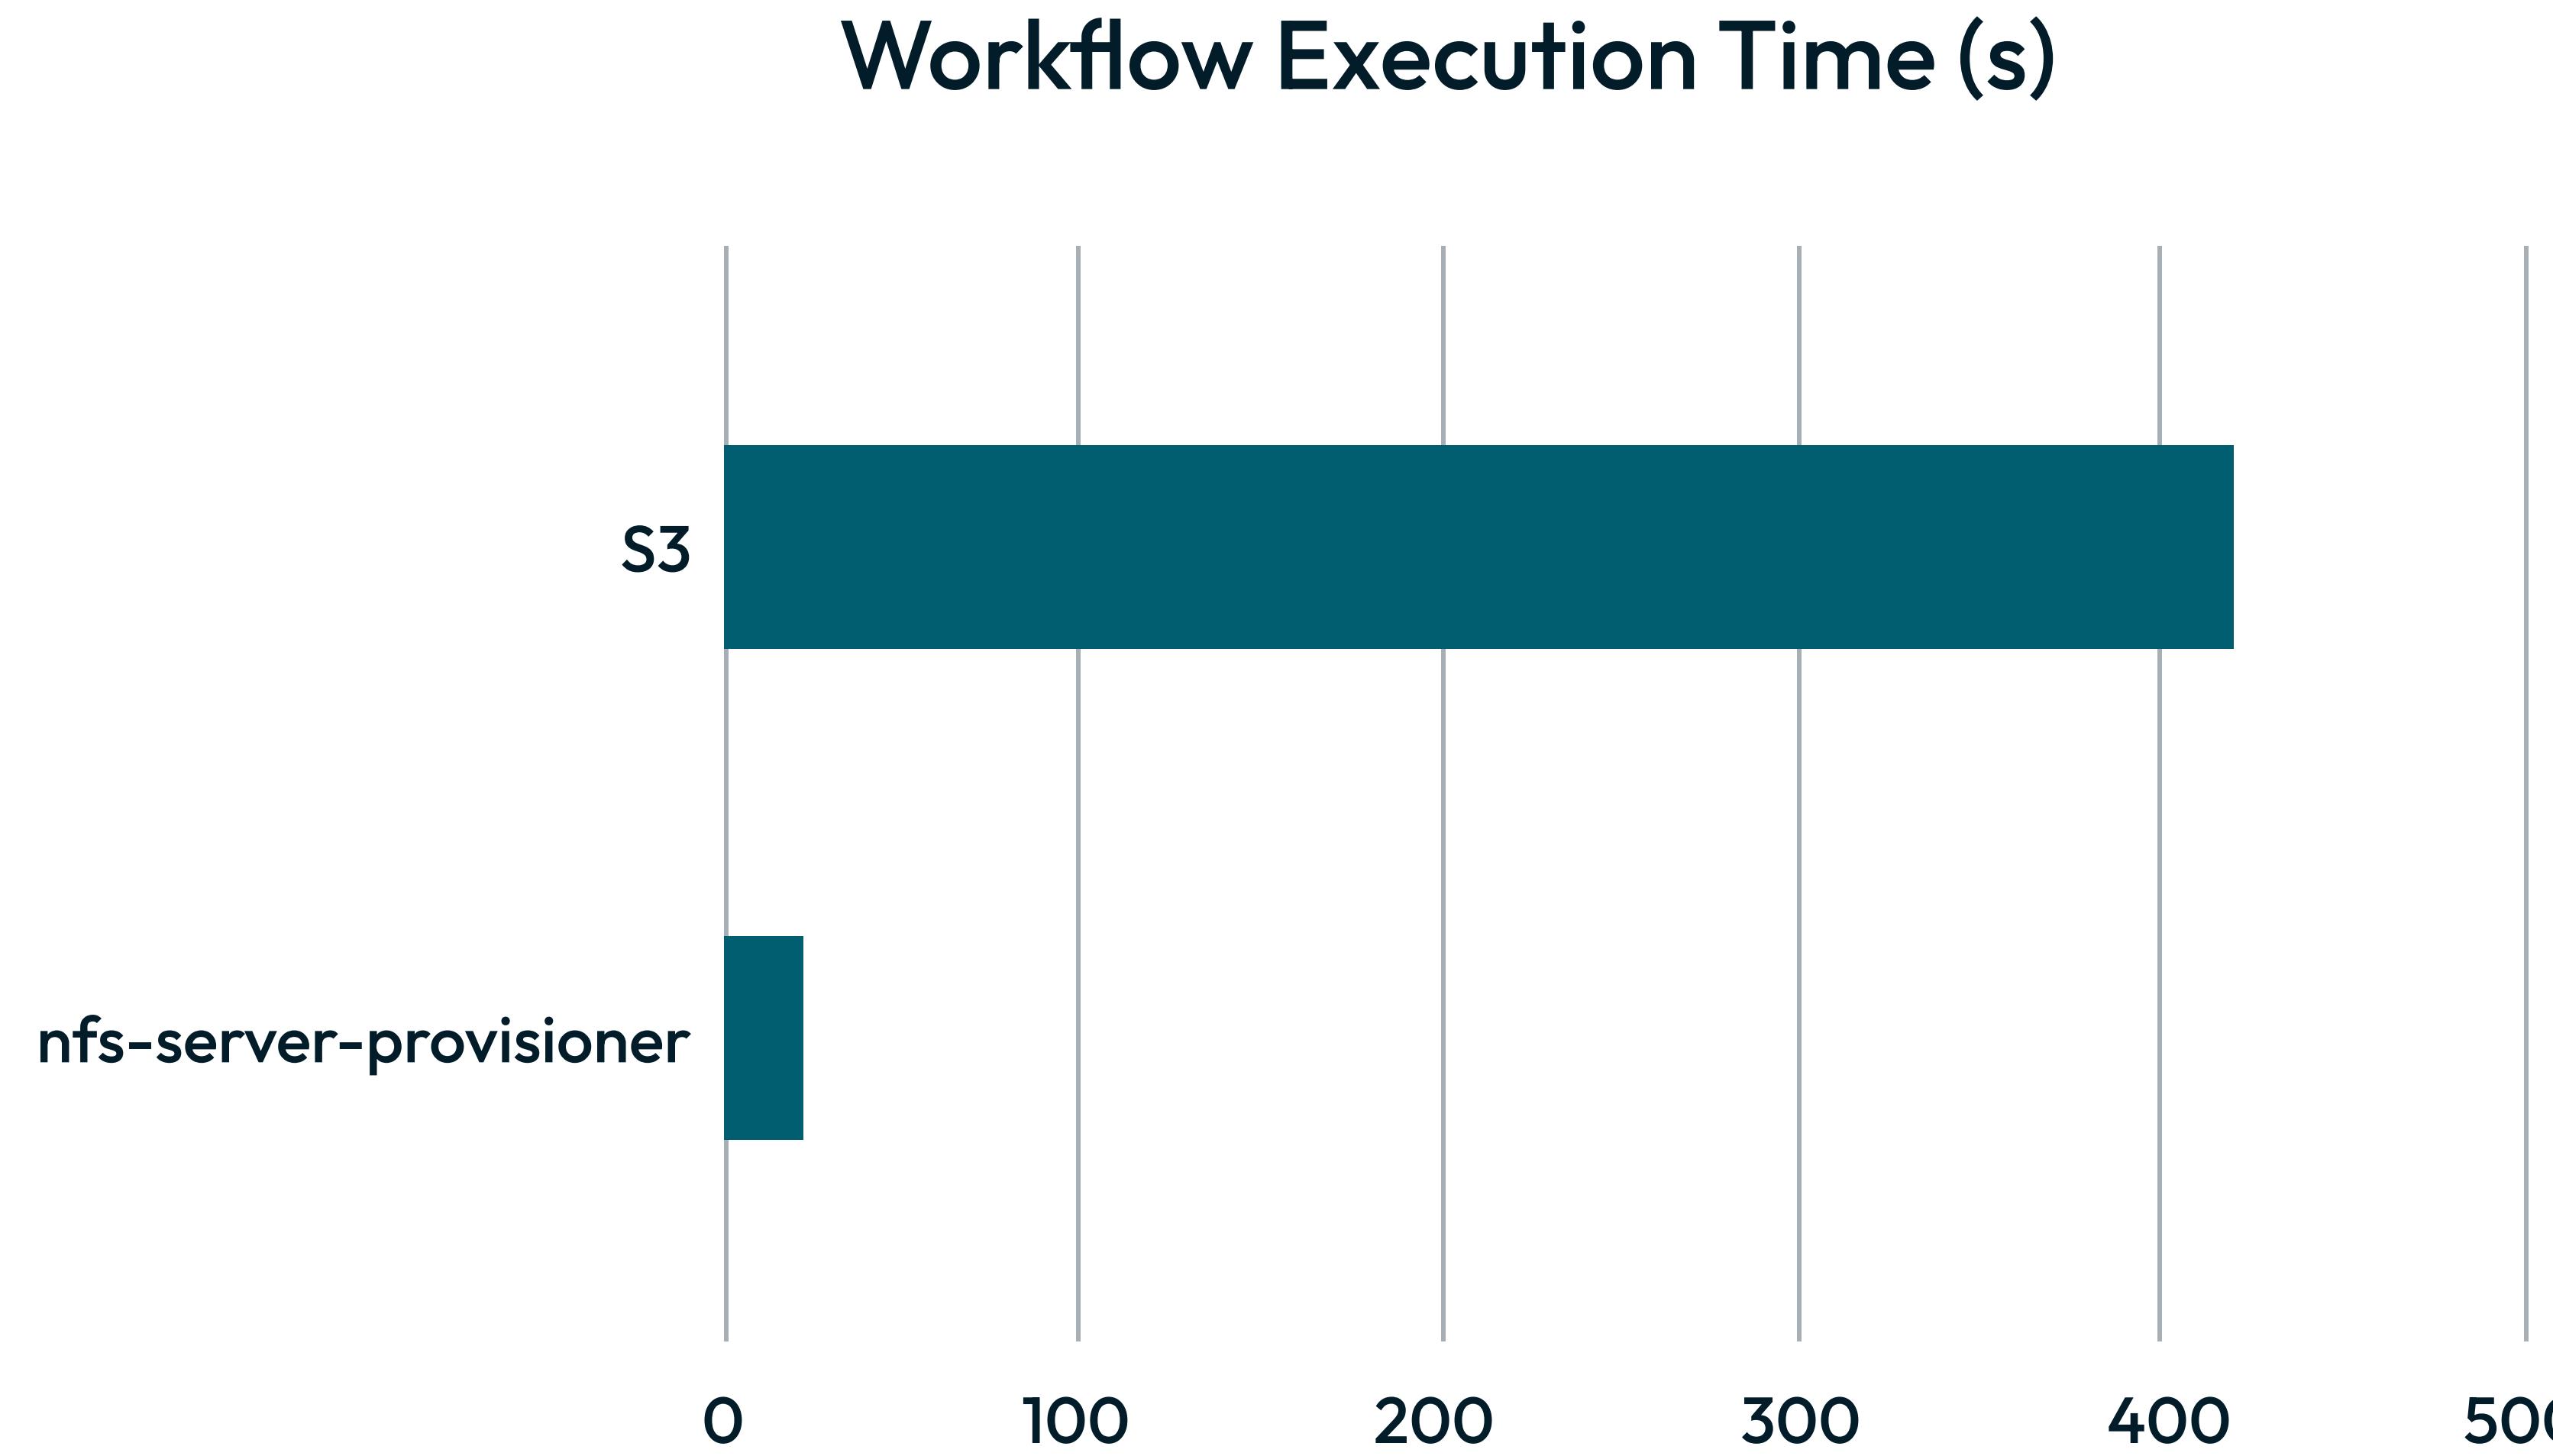


Testing Workflow Speed



These tests can be found at <https://github.com/pipekit/talk-demos>

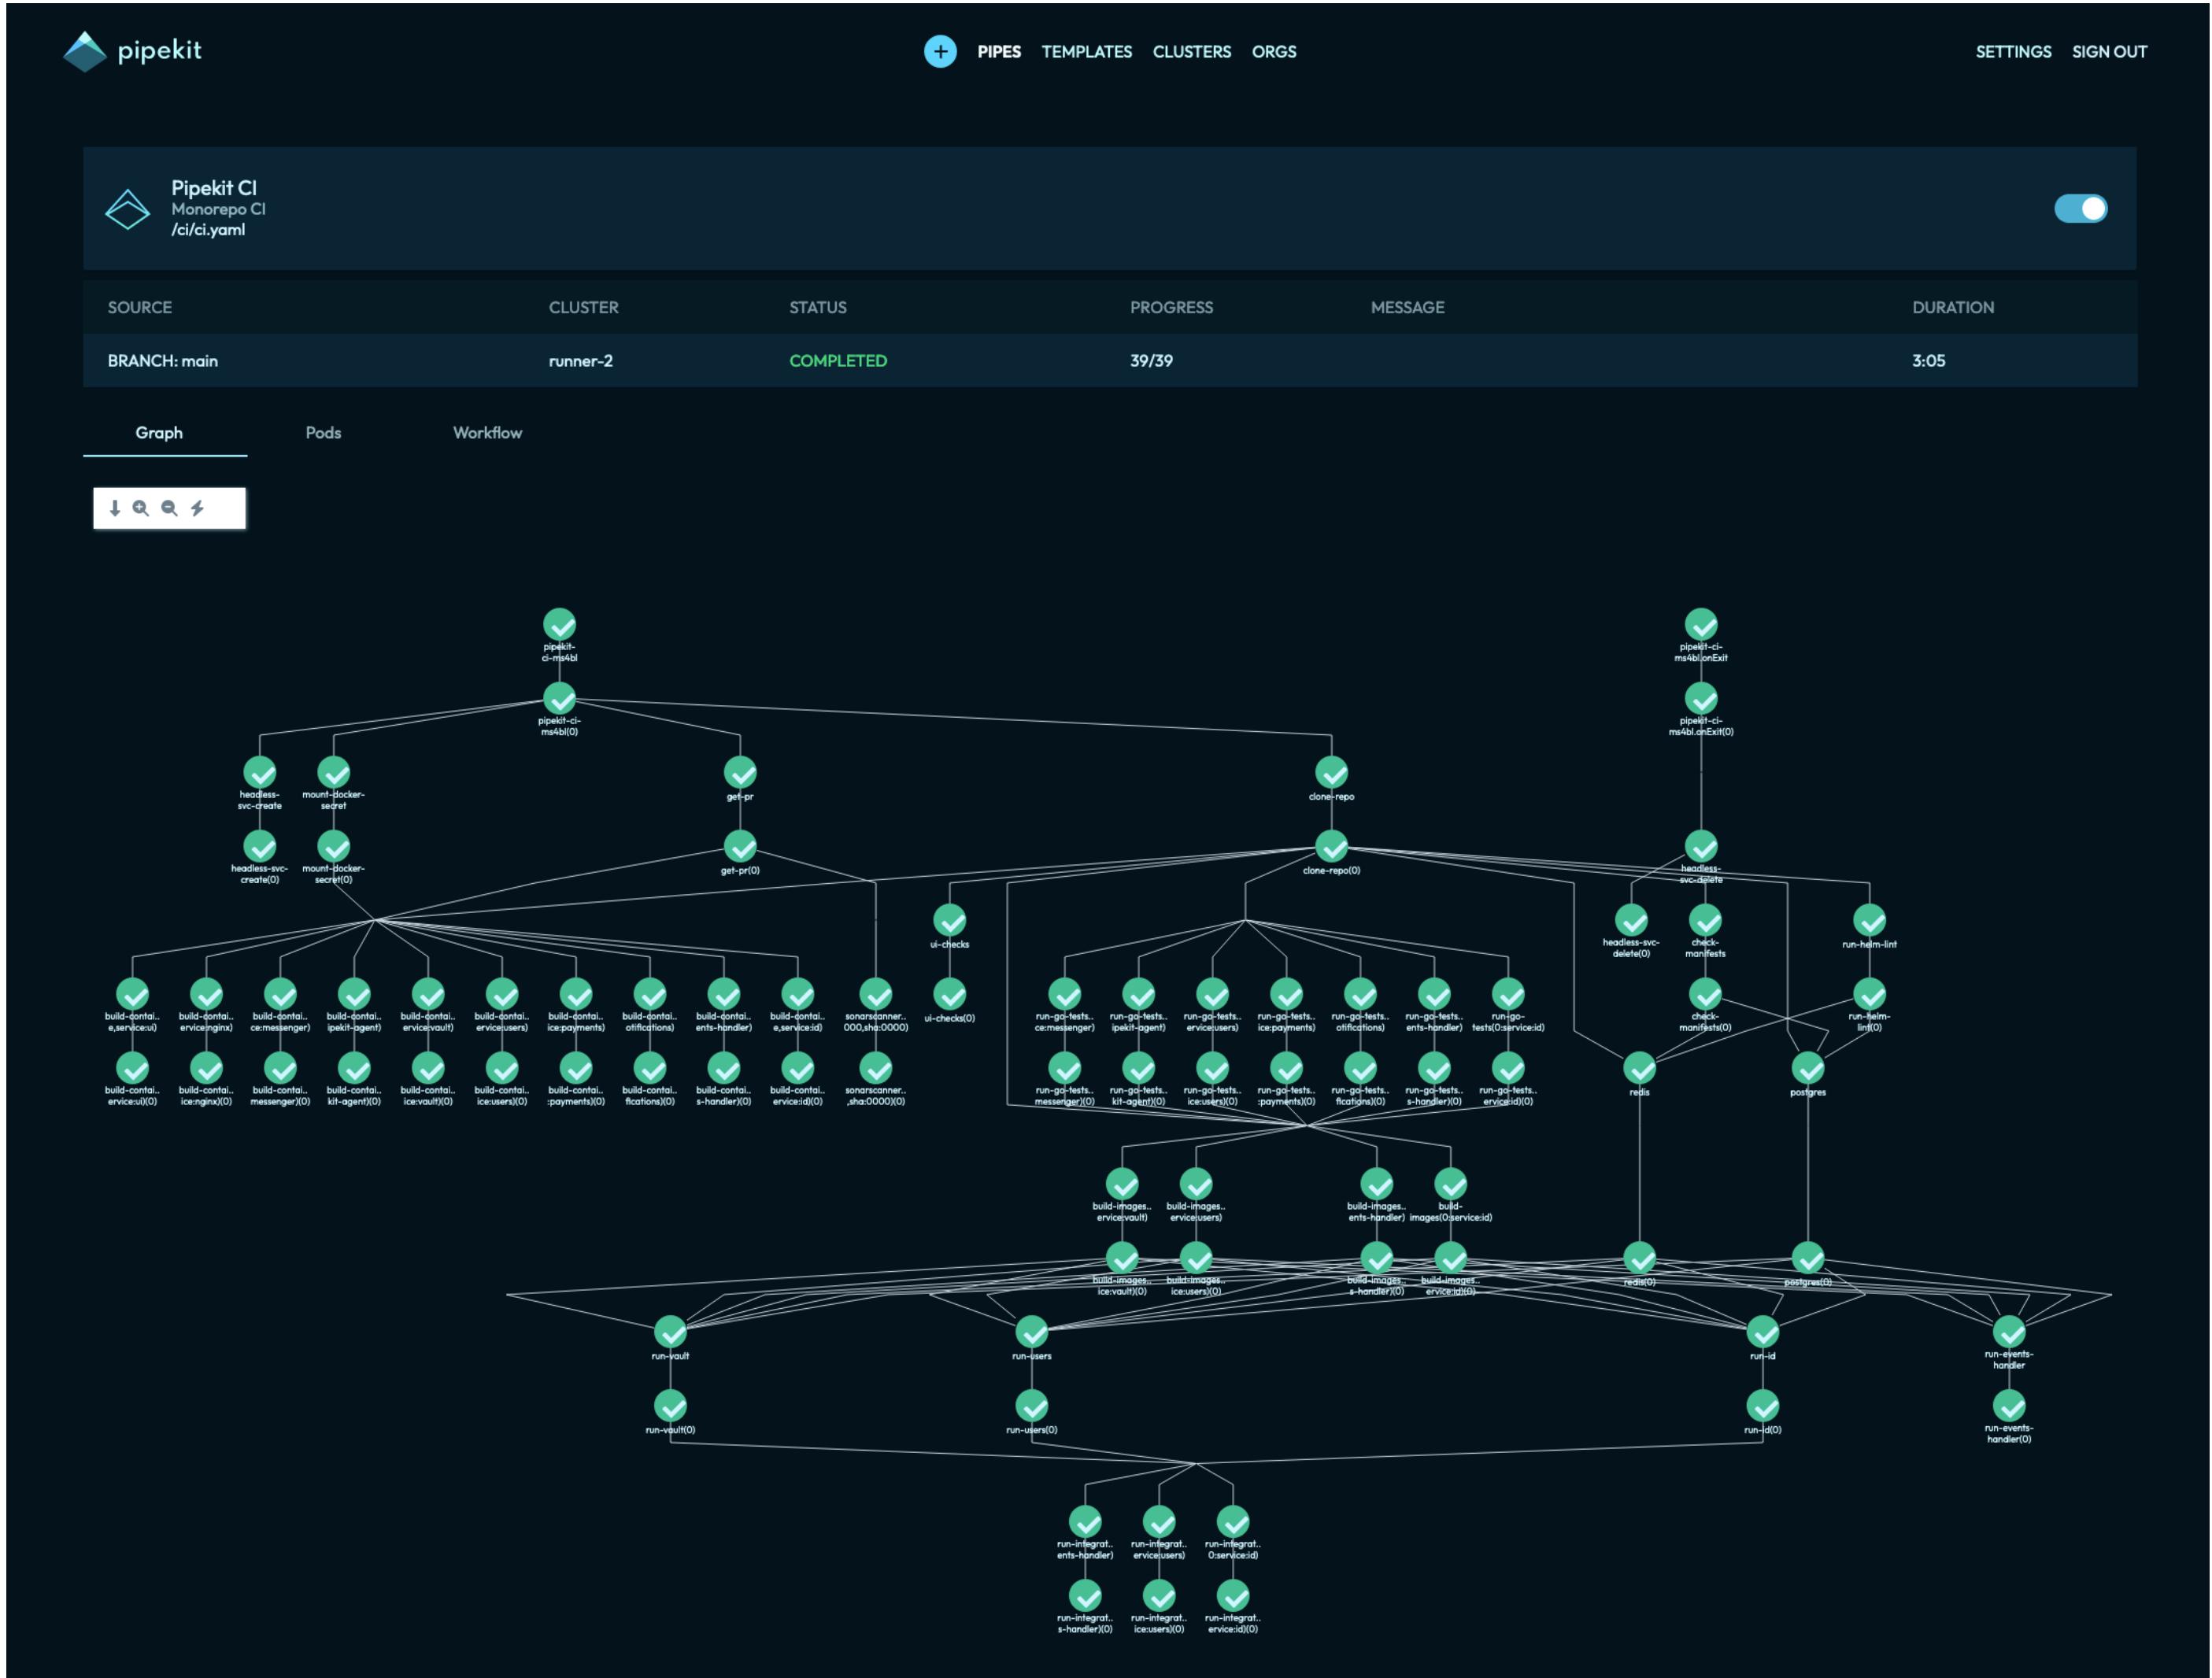
Timing



NFS-Server Provisioner - Tips

- Fatten up your node disk
- Copy from NFS disk to ephemeral at each step if expecting a lot of parallel reads and writes.
- Install NFS-Server-Provisioner twice
 - One backed by an EBS disk for semi-persistent data
 - One backed by ephemeral node disks, for transient data

Conclusion



Thank you!



ArgoCon

#ArgoCon2023