



Demystifying Argo Workflows: **An Architectural Deep Dive**



Darko Janjić
*Senior Software Engineer,
Pipekit*



Becky Pauley
*Solutions Engineer,
Venafi Jetstack Consult*



Image credit: Phil Botha, Unsplash

What is Argo Workflows?



Argo Workflows

Kubernetes-native workflow engine supporting DAG and step-based workflows



Argo CD

Declarative continuous delivery with a fully-loaded UI



Argo Events

Advanced Kubernetes deployment strategies such as Canary and Blue-Green made easy.



Argo Rollouts

Event based dependency management for Kubernetes.

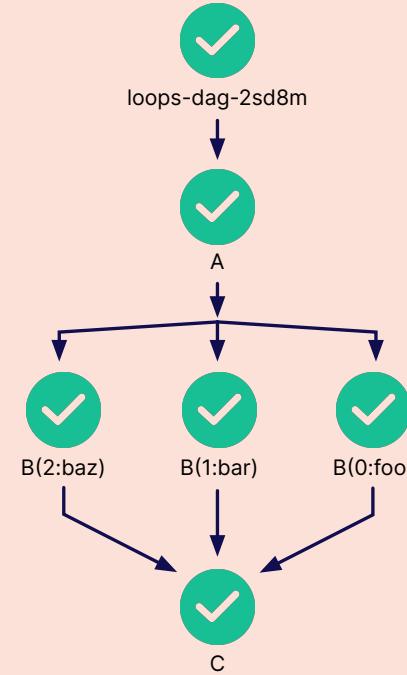
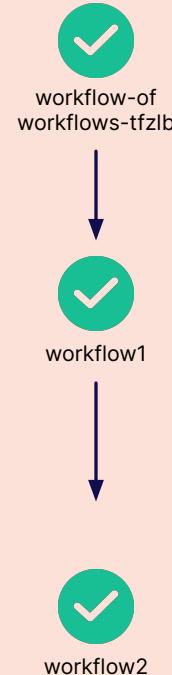
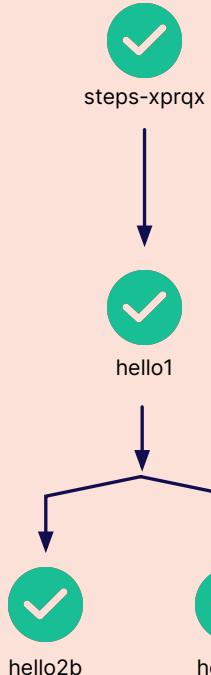
A workflow is...

```
apiVersion: argoproj.io/v1alpha1
kind: Workflow          # new type of k8s spec
metadata:
  generateName: hello-world- # name of the workflow spec
spec:
  entrypoint: whalesay      # invoke the whalesay template
  templates:
    - name: whalesay        # name of the template
      container:
        image: docker/whalesay
        command: [ cowsay ]
        args: [ "hello world" ]
      resources: # limit the resources
        limits:
          memory: 32Mi
          cpu: 100m
```

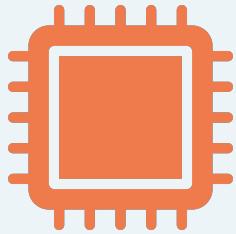


hello-world-bwlq4

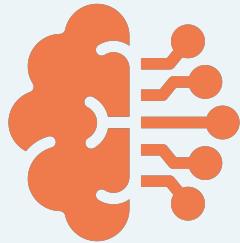
A workflow is...



What is Argo Workflows?



Easily run
compute-intensive jobs



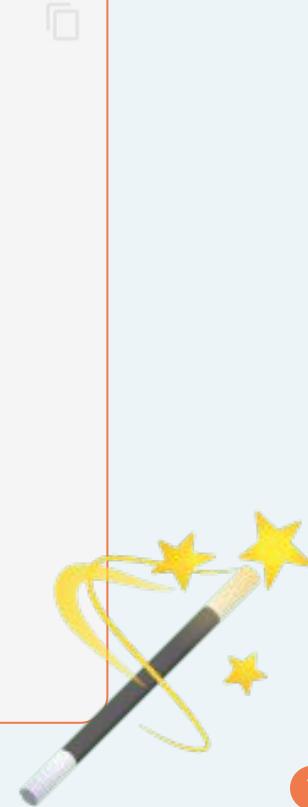
Machine learning,
data processing



Automating infrastructure
and CI/CD pipelines

What is Argo Workflows?

```
apiVersion: argoproj.io/v1alpha1
kind: Workflow          # new type of k8s spec
metadata:
  generateName: hello-world- # name of the workflow spec
spec:
  entrypoint: whalesay      # invoke the whalesay template
  templates:
    - name: whalesay        # name of the template
      container:
        image: docker/whalesay
        command: [ cowsay ]
        args: [ "hello world" ]
      resources: # limit the resources
        limits:
          memory: 32Mi
          cpu: 100m
```



What exists in my cluster?

```

customresourcedefinition.apiextensions.k8s.io/cronworkflows.argoproj.io created
customresourcedefinition.apiextensions.k8s.io/clusterworkflowtemplates.argoproj.io created
customresourcedefinition.apiextensions.k8s.io/workflowartifactctasks.argoproj.io created
customresourcedefinition.apiextensions.k8s.io/workfloweventbindings.argoproj.io
customresourcedefinition.apiextensions.k8s.io/workflowtaskresults.argoproj.io created
customresourcedefinition.apiextensions.k8s.io/workflowtasksets.argoproj.io created
customresourcedefinition.apiextensions.k8s.io/workflowtemplates.argoproj.io created
customresourcedefinition.apiextensions.k8s.io/workflows.argoproj.io created
serviceaccount/argo created
serviceaccount/argo-server created
role.rbac.authorization.k8s.io/argo-role created
clusterrole.rbac.authorization.k8s.io/argo-aggregate-to-admin created
clusterrole.rbac.authorization.k8s.io/argo-aggregate-to-edit created
clusterrole.rbac.authorization.k8s.io/argo-aggregate-to-view created
clusterrole.rbac.authorization.k8s.io/argo-cluster-role created
clusterrole.rbac.authorization.k8s.io/argo-server-cluster-role created
rolebinding.rbac.authorization.k8s.io/argo-binding created
clusterrolebinding.rbac.authorization.k8s.io/argo-binding created
clusterrolebinding.rbac.authorization.k8s.io/argo-server-binding created
configmap/workflow-controller-configmap created
service/argo-server created
priorityclass.scheduling.k8s.io/workflow-controller created
deployment.apps/argo-server created
deployment.apps/workflow-controller created

```



CRDs

Workflow

WorkflowTaskSet

CronWorkflow

WorkflowTaskResult

WorkflowTemplate

WorkflowEventBindings

ClusterWorkflowTemplate

WorkflowArtifactGCTasks



Deployments

Workflow Controller

Argo Server

Custom Resource Definitions

```

1  apiVersion: apiextensions.k8s.io/v1
2  kind: CustomResourceDefinition
3  metadata:
4    name: workflows.argoproj.io
5  spec:
6    group: argoproj.io
7    names:
8      kind: Workflow
9      listKind: WorkflowList
10     plural: workflows
11     shortNames:
12       - wf
13     singular: workflow
14   scope: Namespaced
15   versions:
16     - additionalPrinterColumns:
17       - description: Status of the workflow
18         jsonPath: .status.phase
19         name: Status
20         type: string
21       - description: When the workflow was started
22         format: date-time
23         jsonPath: .status.startedAt
24         name: Age
25         type: date
26       - description: Human readable message indicating a condition.
27         is in this condition.
28         jsonPath: .status.message
29         name: Message
30         type: string

```



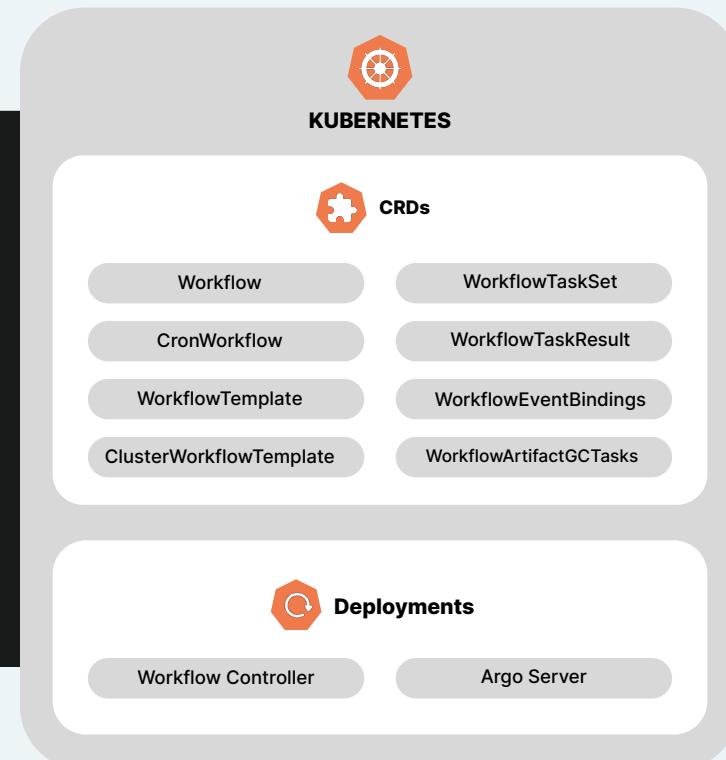
```

apiVersion: argoproj.io/v1alpha1
kind: Workflow
metadata:
  generateName: hello-world-
spec:
  entrypoint: whalesay
  templates:
    - name: whalesay
      container:
        image: docker/whalesay
        command: [ cowsay ]
        args: [ "hello world" ]
        resources: # limit the resources
          limits:
            memory: 32Mi
            cpu: 100m

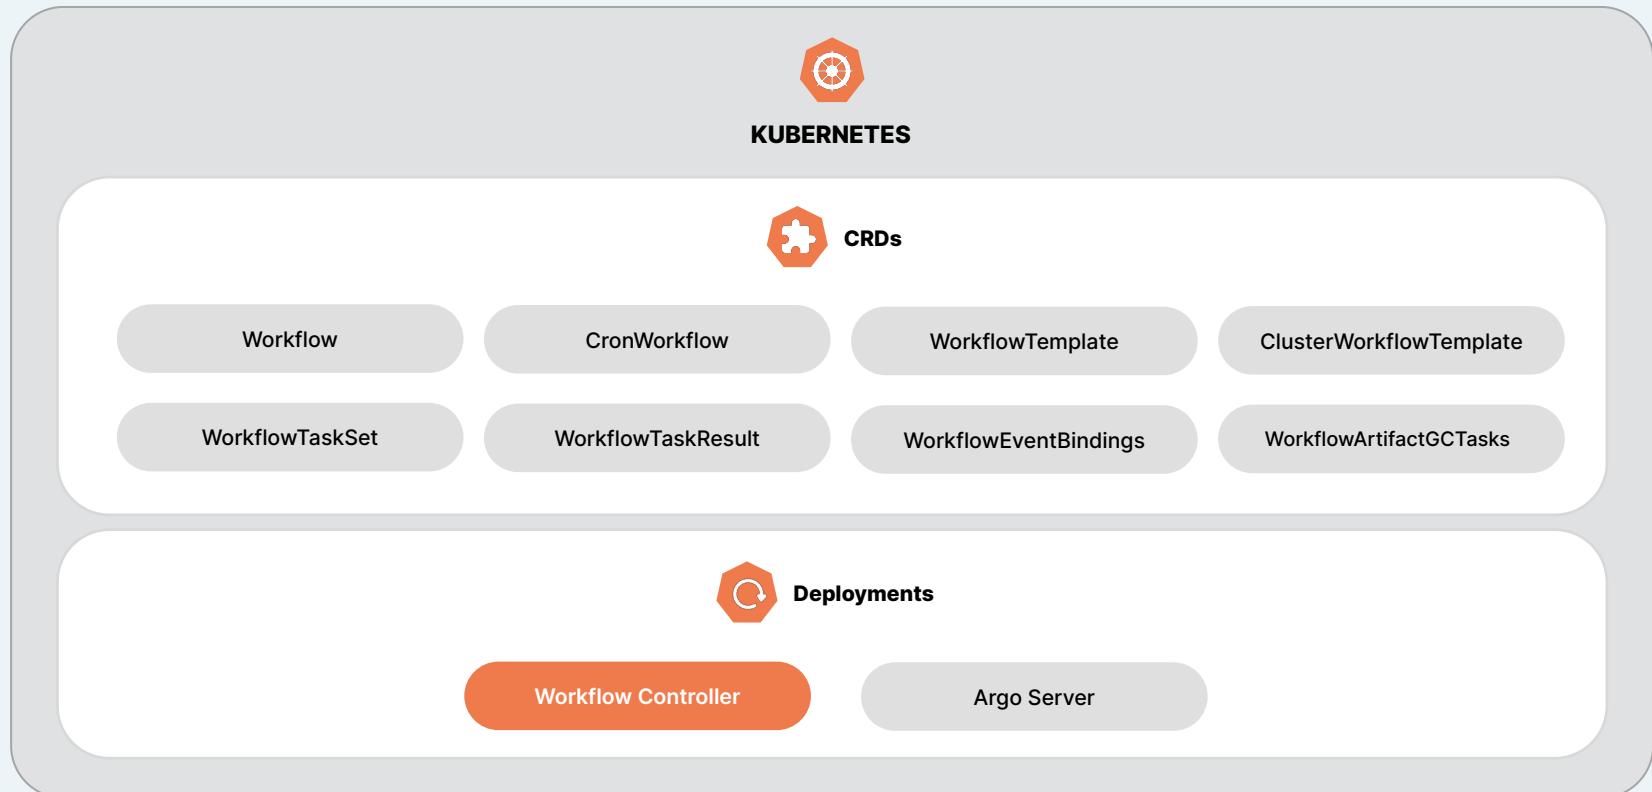
```

Custom Resource Definitions

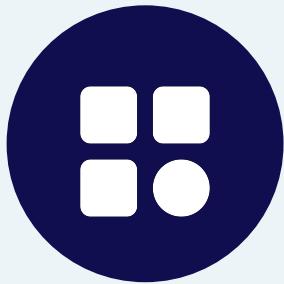
```
NAME
clusterworkflowtemplates.argoproj.io
cronworkflows.argoproj.io
workflowartifactgctasks.argoproj.io
workfloweventbindings.argoproj.io
workflows.argoproj.io
workflowtaskresults.argoproj.io
workflowtasksets.argoproj.io
workflowtemplates.argoproj.io
```



How do we orchestrate Workflows?



Workflow Controller



Main component of
Argo Workflows

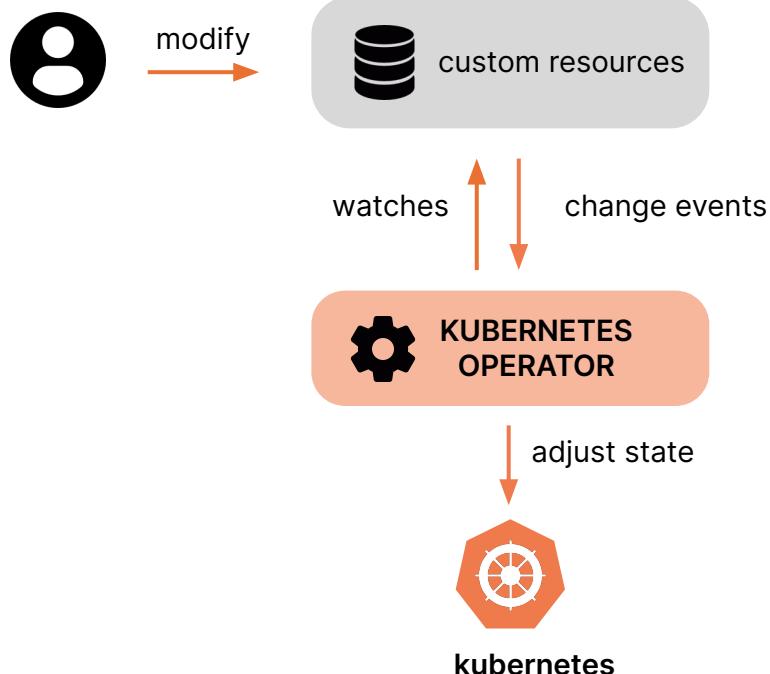


Works with
custom resources



Manages workflows and
everything Argo Workflows related
(CronWorkflows,
WorkflowTemplates...)

Workflow Controller



- Implements Kubernetes operator pattern
- Mostly talks with Kubernetes
- Can run in HA or have namespaced deployment

Workflow Controller - modes

HIGH AVAILABILITY DEPLOYMENT



Workflow
Controller



Workflow
Controller



Workflow
Controller



Namespace A



Workflow Controller - modes

NAMESPACEDEPLOYMENT



Workflow Controller



Namespace A



Workflow Controller



Namespace B



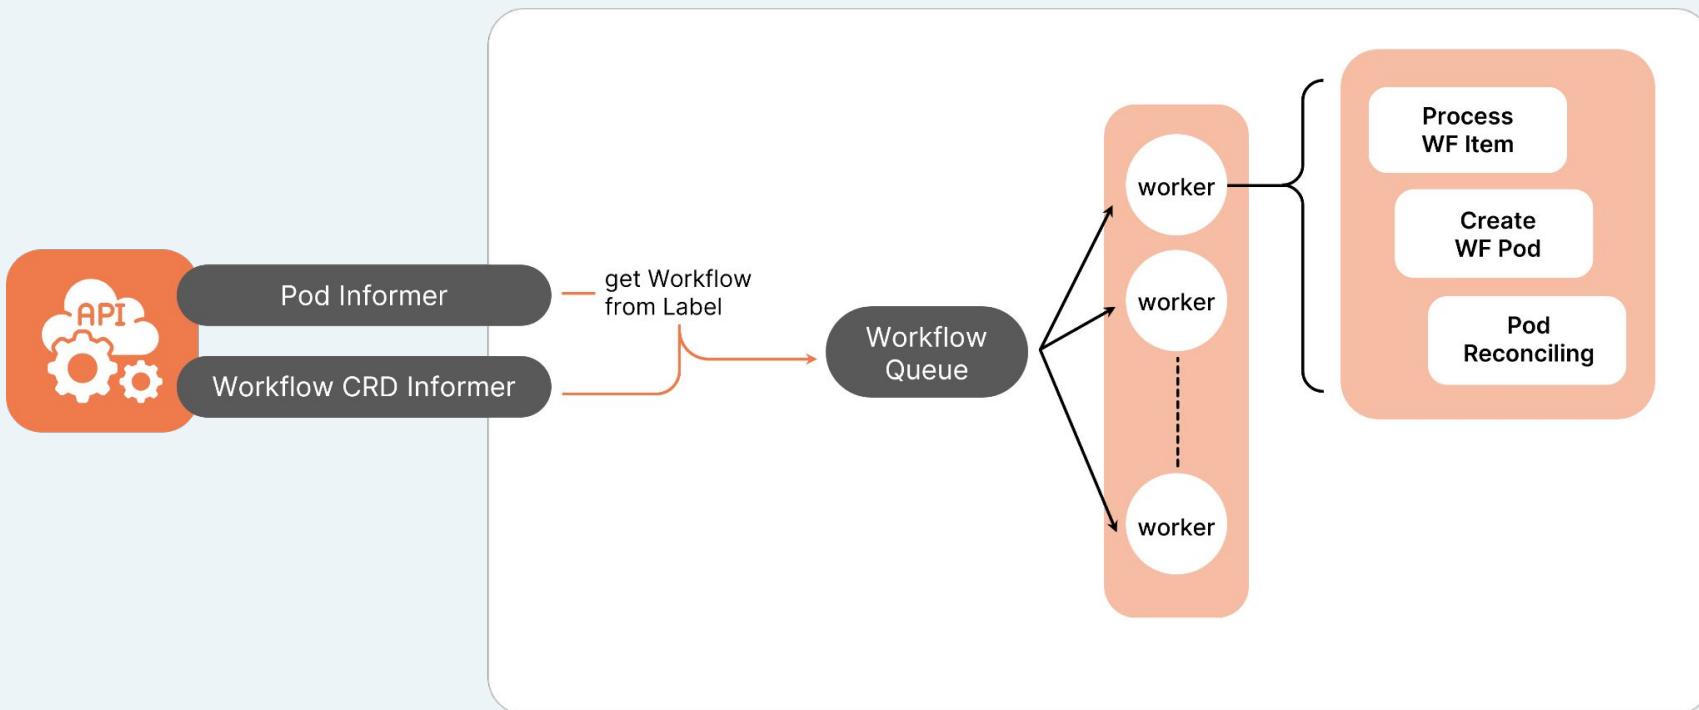
Workflow Controller



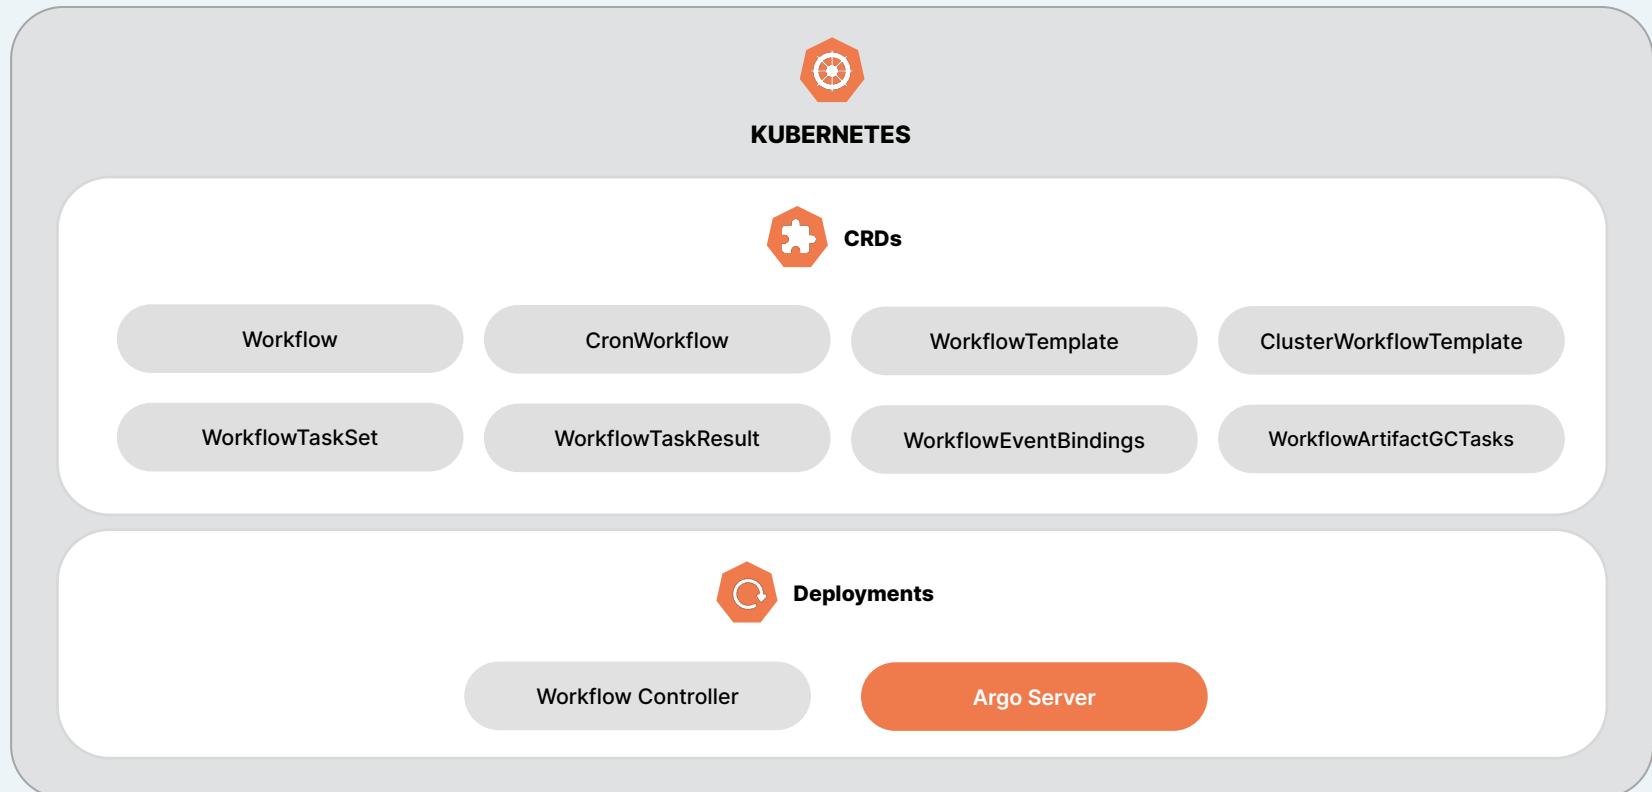
Namespace C



Workflow Controller - modes



How do we orchestrate Workflows?



Argo Server



Mostly used for the communication with the outside world



For people that are not Kubernetes experts



Exposes Argo Workflows API



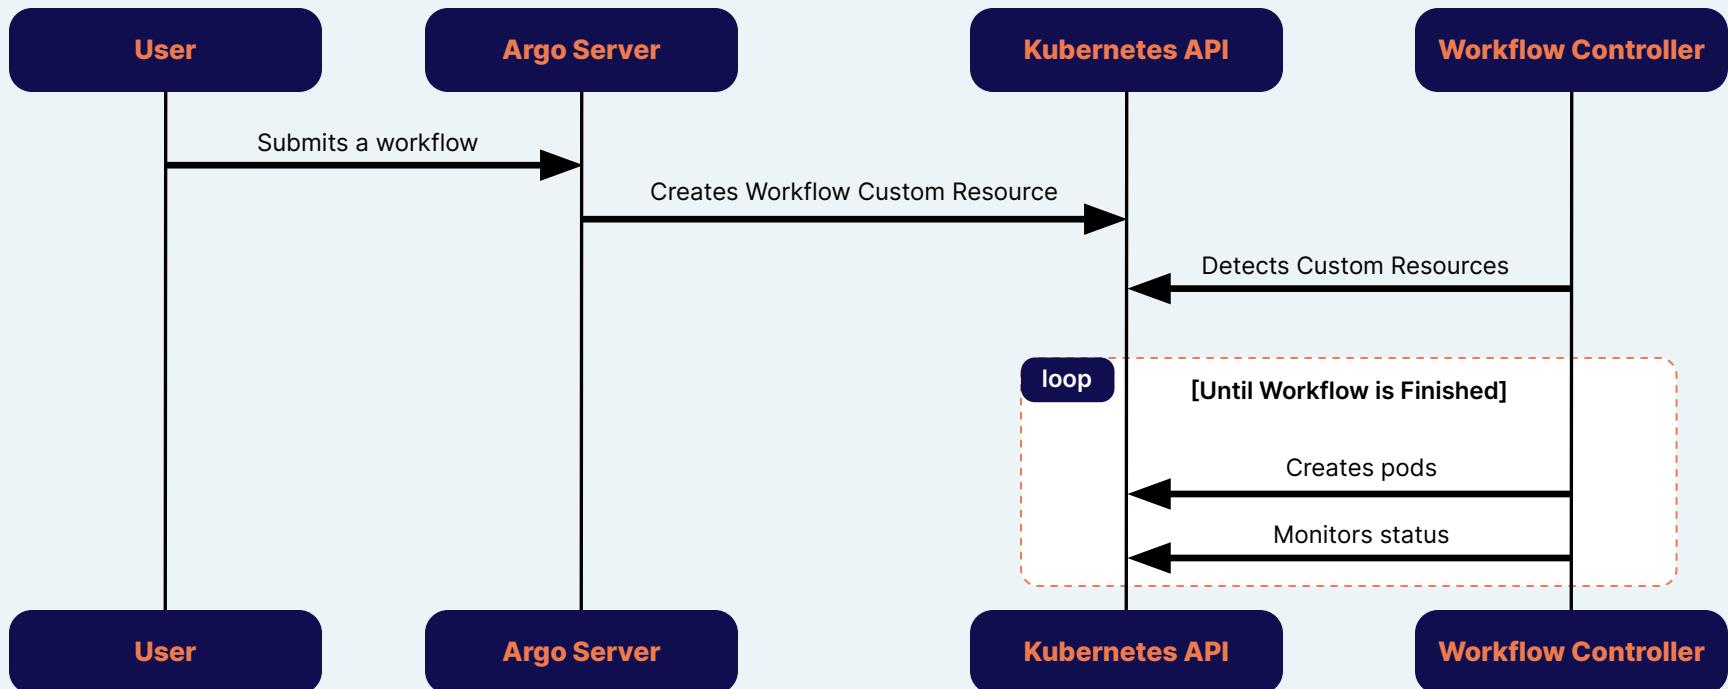
Provides UI

Argo Server

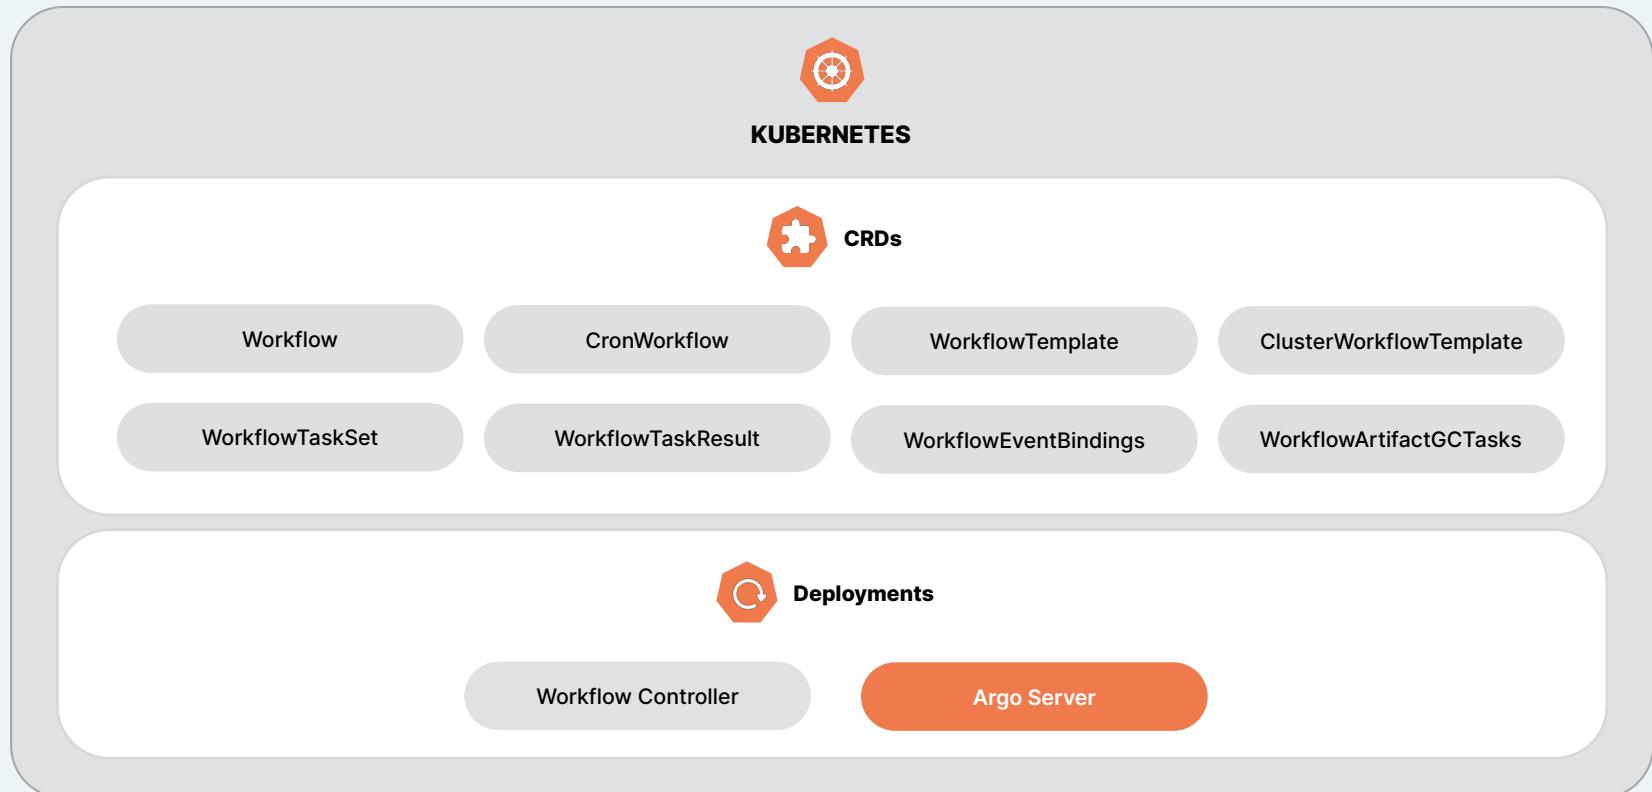


- Argo Workflow CLI is using Argo Server
- Not a mission-critical component
- Does the authentication
- Interacts with Kubernetes API
- Features like Workflow Archive, Offloading large workflows and Events need Argo Server

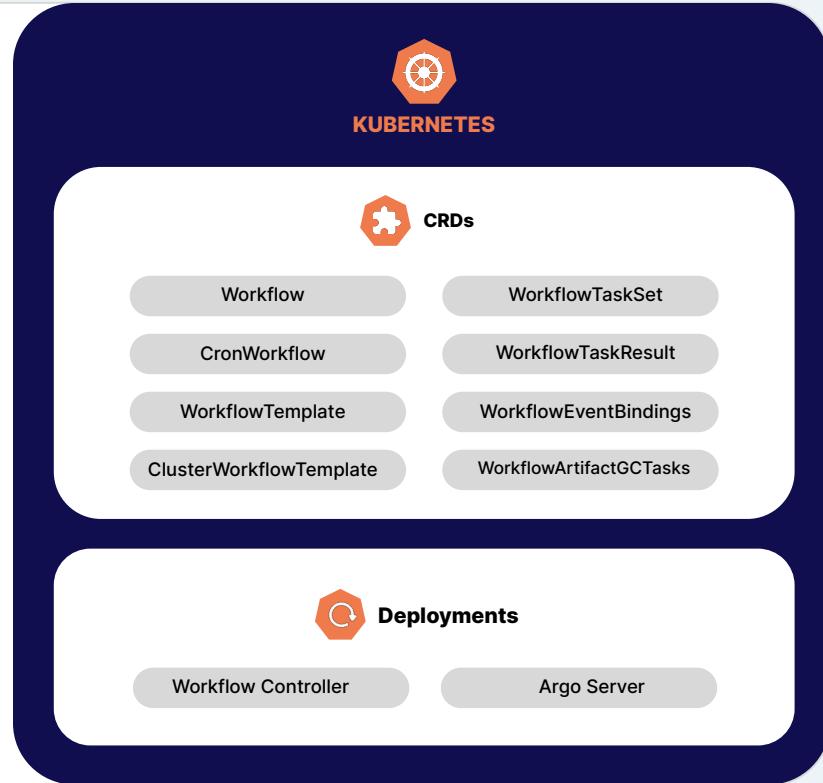
The Flow



Revisiting the Big Picture



The humble Pod

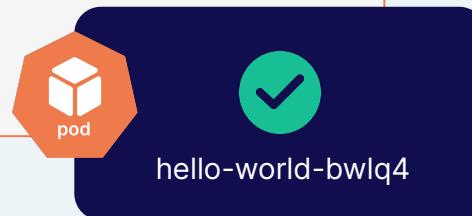


The basis of our Workflow steps

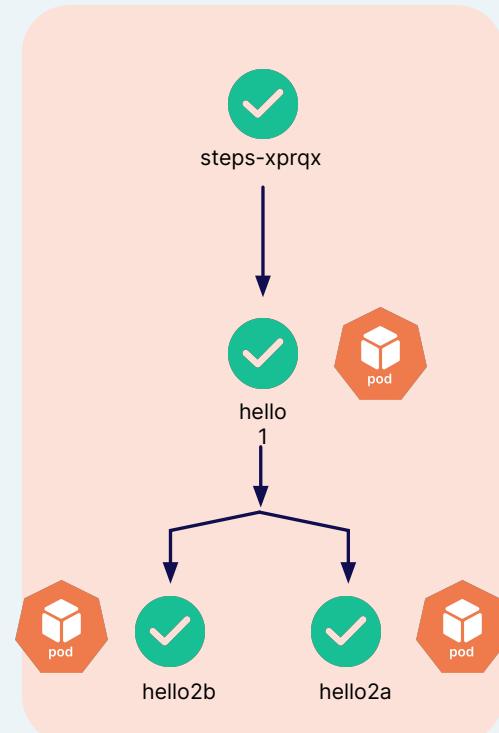
```

apiVersion: argoproj.io/v1alpha1
kind: Workflow          # new type of k8s spec
metadata:
  generateName: hello-world- # name of the workflow spec
spec:
  entrypoint: whalesay      # invoke the whalesay template
  templates:
    - name: whalesay        # name of the template
      container:
        image: docker/whalesay
        command: [ cowsay ]
        args: [ "hello world" ]
        resources: # limit the resources
          limits:
            memory: 32Mi
            cpu: 100m

```

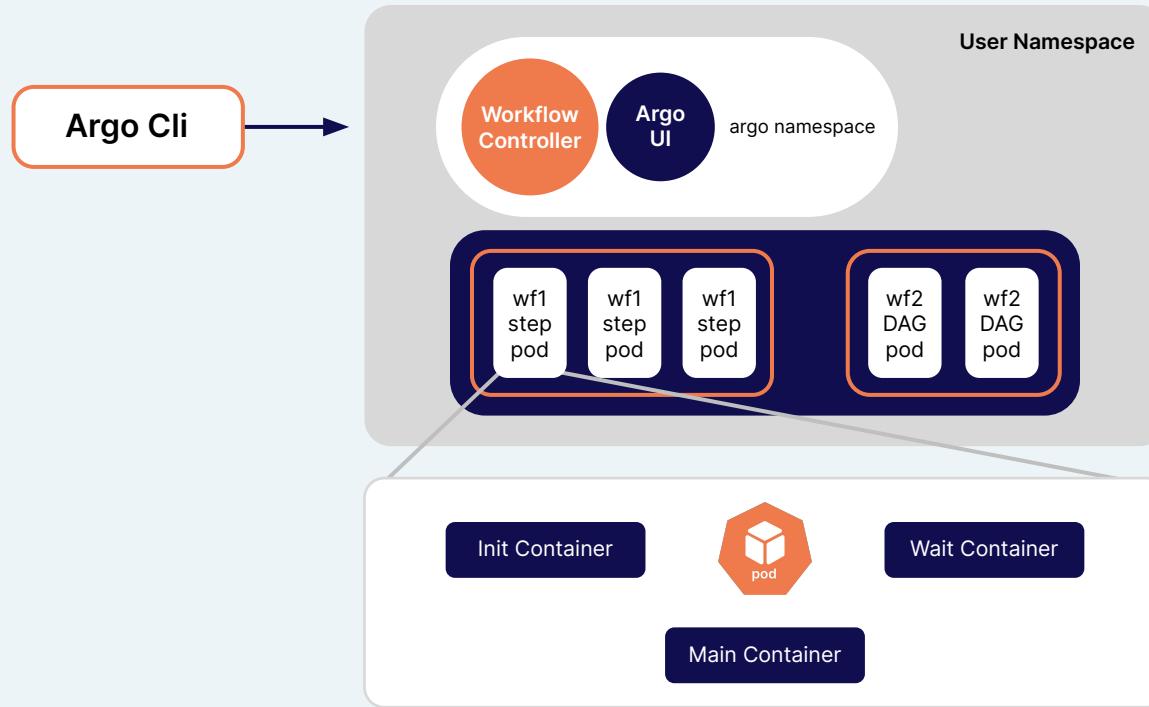


✗ suspend, workflow of workflows



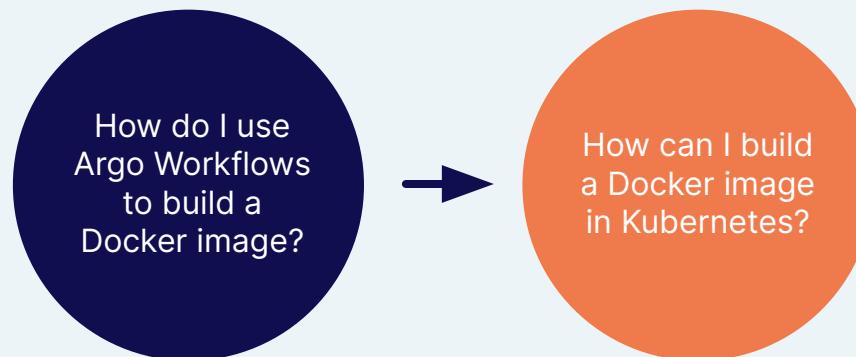
Step == Pod

ARGO Workflow overview



How do I...?

"If you can do a thing in a normal kubernetes pod, you can do it in a workflow"

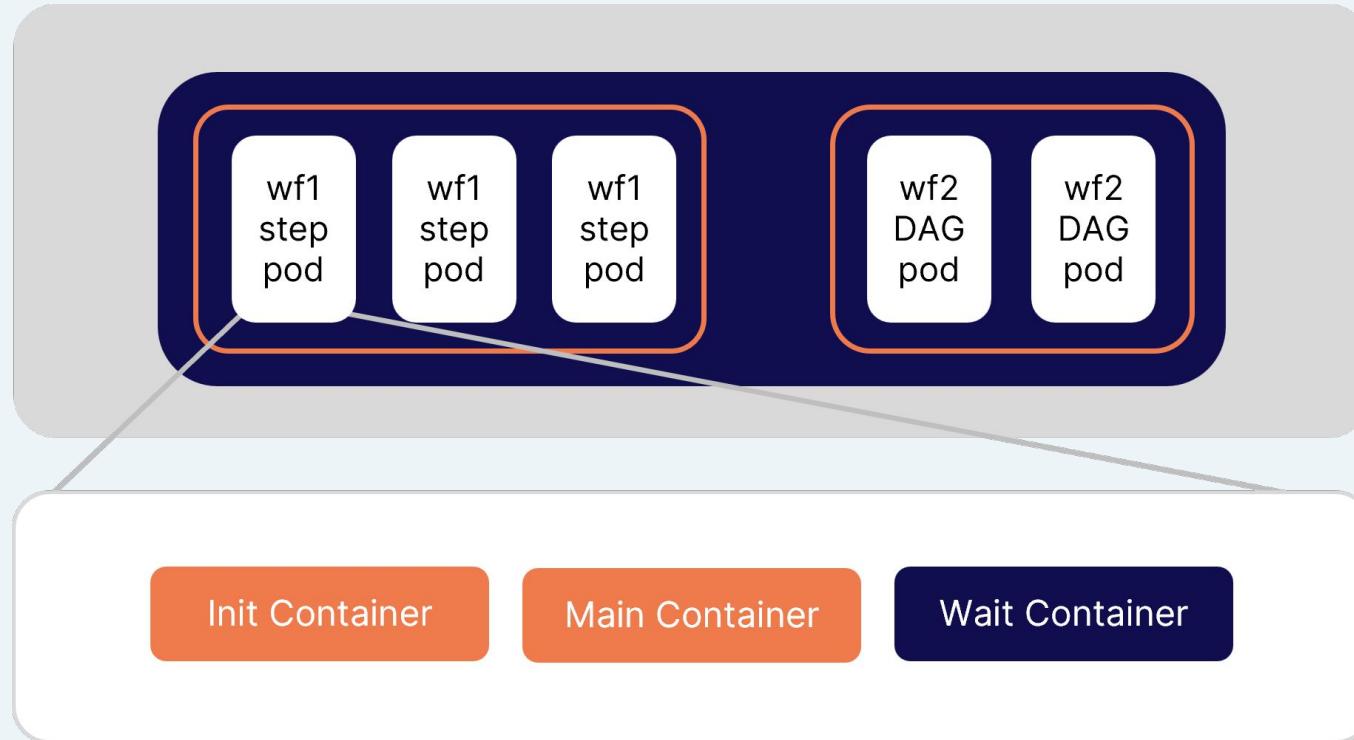


'Podness' helps us understand...

- Resource utilisation
- RBAC
- Mount volumes/other containers
- Workspace
- Very short-lived steps
- Metrics



The anatomy of a Step (a Pod!)



We can see this when...

```
spec:
  entrypoint: python-script-example
  templates:
    - name: python-script-example
      steps:
        - - name: generate
          template: gen-random-int
        - - name: print
          template: print-message
          arguments:
            parameters:
              - name: message
                value: "{{steps.generate.outputs.result}}"

    - name: gen-random-int
      script:
        image: python:alpine3.6
        command: [python]
        source: |
          import random
          i = random.randint(1, 100)
          print(i)
```

```
apiVersion: argoproj.io/v1alpha1
kind: Workflow
metadata:
  generateName: container-set-template-
spec:
  entrypoint: main
  templates:
    - name: main
      volumes:
        - name: workspace
          emptyDir: { }
      containerSet:
        volumeMounts:
          - mountPath: /workspace
            name: workspace
        containers:
          - name: a
            image: argoproj/argosay:v2
            command: [sh, -c]
            args: ["echo 'a: hello world' >> /workspace/message"]
          - name: b
            image: argoproj/argosay:v2
            command: [sh, -c]
            args: ["echo 'b: hello world' >> /workspace/message"]
```

But what about...

```

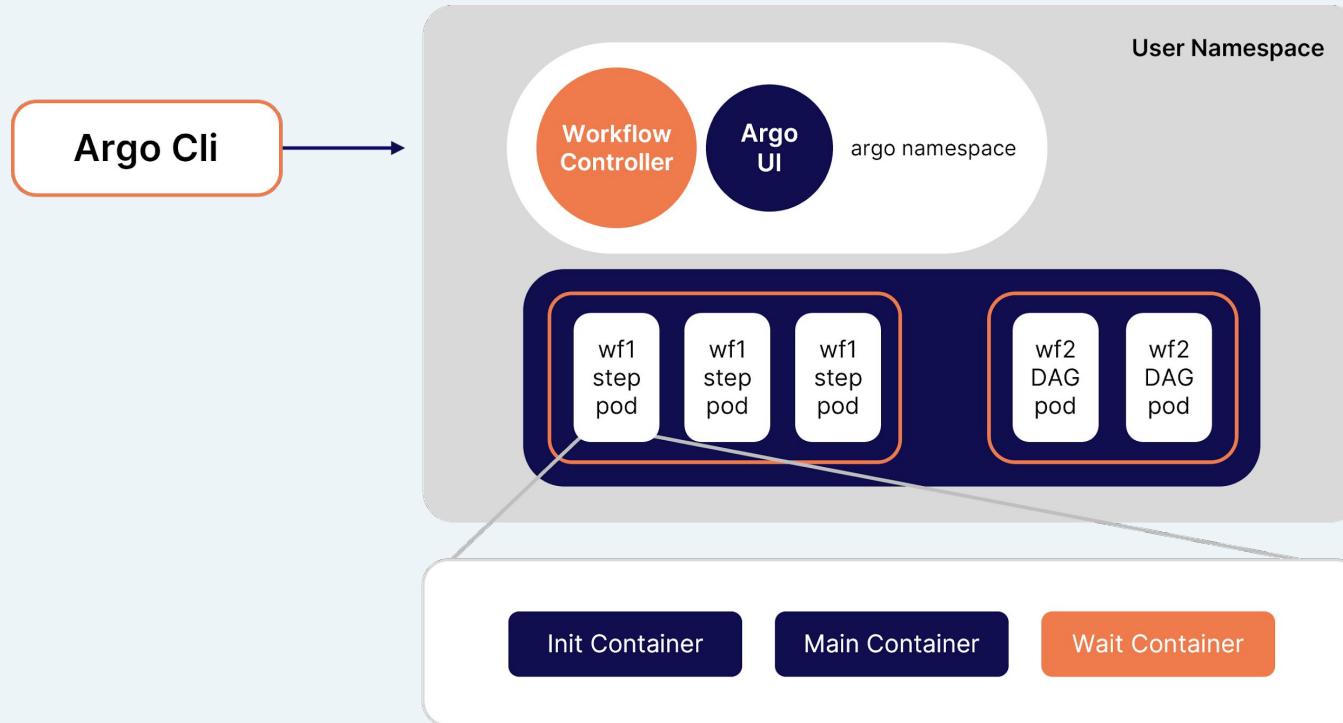
entrypoint: main
templates:
- name: main
  steps:
    - - name: get-google-homepage
      template: http
      arguments:
        parameters: [{name: url, value: "https://www.google.com"}]
    - name: http
      inputs:
        parameters:
          - name: url
      http:
        timeoutSeconds: 20 # Default 30
        url: "{{inputs.parameters.url}}"
        method: "GET" # Default GET
        headers:
          - name: "x-header-name"
            value: "test-value"
    # Template will succeed if evaluated to true, otherwise will fail
    # Available variables:
    # request.body: string, the request body
    # request.headers: map[string][]string, the request headers
    # response.url: string, the request url
    # response.method: string, the request method
    # response.statusCode: int, the response status code
    # response.body: string, the response body
    # response.headers: map[string][]string, the response headers
    successCondition: "response.body contains \"google\"" # available since v3
    body: "test body" # Change request body
  
```

```

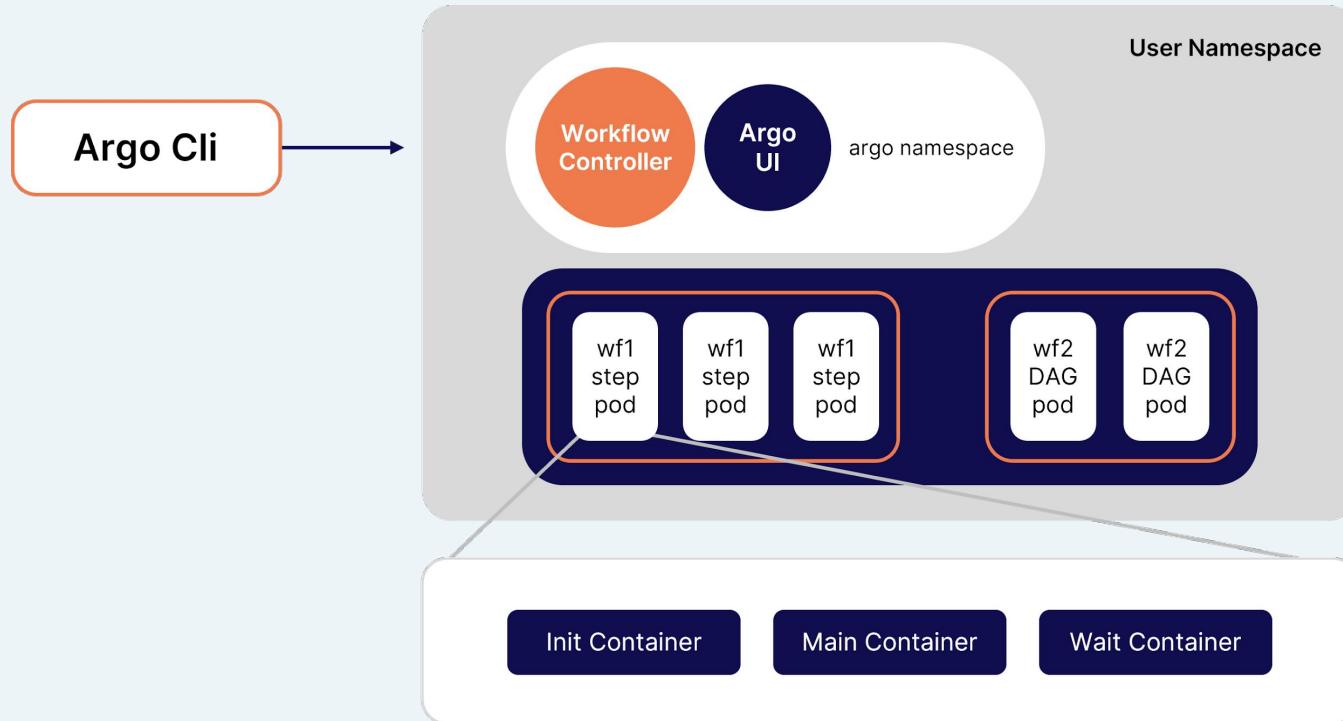
- name: create-configmap
  resource:
    action: create
    manifest: |
      apiVersion: v1
      kind: ConfigMap
      metadata:
        name: resource-delete-with-flags
        labels:
          cleanup: "true"
      data:
        key: value
  
```



The anatomy of a Step (a Pod!)



The Big Picture



Want to learn more?



Try the Killer Coda course
<https://killercoda.com/argoproj/course/argo-workflows/>



Read through the docs
<https://argo-workflows.readthedocs.io>



Visit GitHub (especially the examples directory)
<https://github.com/argoproj/argo-workflows>



Stay up to date with our blogs
<https://pipekit.io/blog>
<https://venafi.com/blog/category/cloud-native/>



Visit Pipekit Office Hours

About Pipekit



Argo experts & maintainers



Save engineering time and up to 60% on compute costs



Add a team of **3 Argo maintainers & 5 Argo contributors** to your Slack



Serving **startups & Fortune 500 enterprises** since 2021:

- **Enterprise Support for Argo** → Ideal for Platform Eng teams scaling with Argo
- **Control Plane for Argo Workflows** → Ideal for Data teams who don't speak K8s OR multi-cluster setups

Cloud native ecosystem knowledge and expertise

Venafi Jetstack Consult provides strategic consulting and advisory services for companies that are embracing cloud native



<https://venafi.com/jetstack-consult/>



Thank you!



Darko Janjić
Senior Software Engineer, Pipekit

darko@pipekit.io



Becky Pauley
*Solutions Engineer, Venafi Jetstack
Consult*

becky.pauley@venafi.com

Feedback



Q&A



Do you have any questions?