



ArgoCon

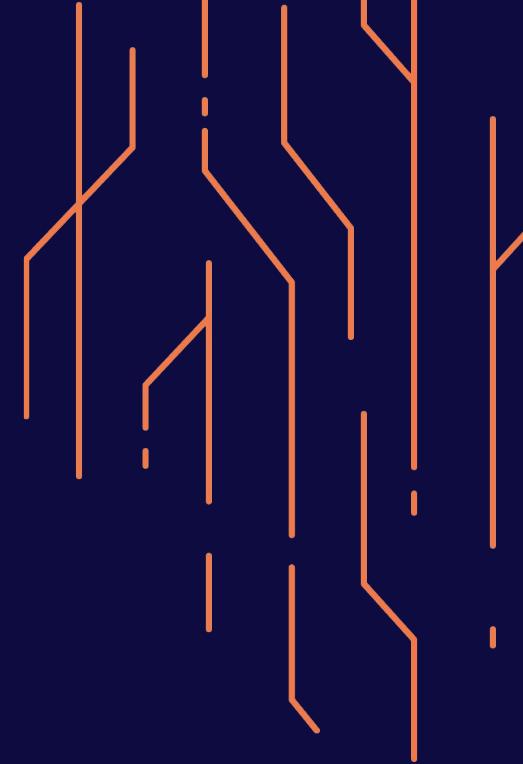
EUROPE



Orchestrating Python Functions Natively with Hera

ArgoCon EU

March 19th, 2024



Speakers



ELLIOT GUNTON



JP ZIVALICH

Bloomberg

Engineering

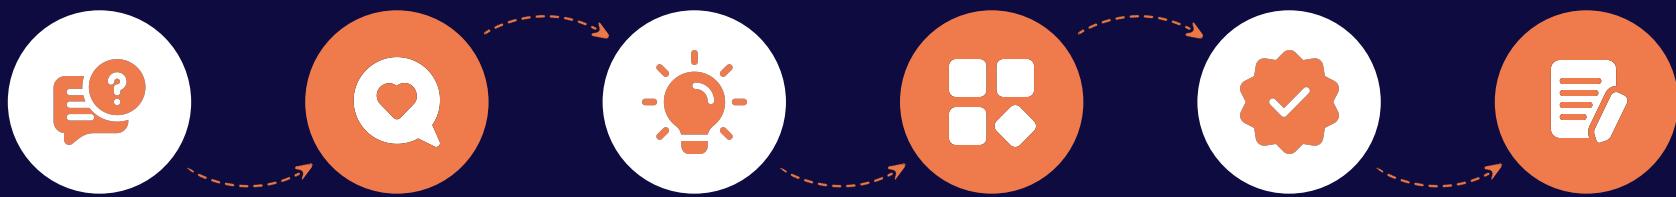
Senior Software Engineer
Hera maintainer



pipekit

CTO & Founder
Hera evangelist

Outline



Why build
on Argo
Workflows?

The Argo
developer
experience

Introducing
Hera

Hera
features and
examples

Case Studies:
How Hera helps
solve problems

Takeaways

The Whats and Whys

Why build on Argo Workflows?

- 1 The de facto Kubernetes Workflow Orchestration Standard
- 2 Allows users to scale more than non-cloud-native tools
- 3 Kubernetes and container-native
- 4 Active and growing community
- 5 Long-term support
- 6 Vendor-neutral



What problems does Argo Workflows have?

YAML is a barrier to entry



Hard to **test**



Hard to **reuse**
(outside of
WorkflowTemplates)



Hard to **maintain** long
Workflows and
WorkflowTemplates

```
apiVersion: argoproj.io/v1alpha1
kind: Workflow
metadata:
  generateName: fine-tune-llm-
spec:
  entrypoint: fine-tune
  onExit: exit
  templates:
    - dag:
        tasks:
          - name: create-ssd-storage-class
            template: create-ssd-storage-class
          - depends: create-ssd-storage-class
            name: create-etcd-stateful-set
            template: create-etcd-stateful-set
          - depends: create-ssd-storage-class
            name: create-etcd-load-balancer
            template: create-etcd-load-balancer
  arguments:
    ...
  
```

Why build on Python?

- ✓ Mature developer experience
- ✓ Extensive libraries
- ✓ Preferred scripting language for many
- ✓ Code completion features
- ✓ Access to testing frameworks
 - ✓ Unit testing of individual functions
 - ✓ System testing of the whole workflow
- ✓ Packaging and versioning



How can we get the best of both?



PYTHON

- Developer experience
- Code completion
- Testing frameworks
- Extensive libraries
- Preferred scripting language
- Packaging and versioning

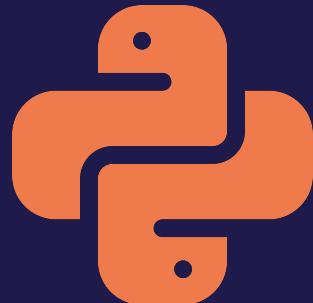


ARGO

- The de facto Kubernetes Workflow Orchestration Standard
- Kubernetes and container-native
- Long-term support
- Active and growing community
- Vendor-neutral

Introducing Hera!

The Python SDK for Argo Workflows



Write your templates as functions,
orchestrate and test – **all in Python!**

Lets you interact with Argo
Workflows entirely through Python!



Did we mention we love Python?!?

Seamless Developer Experience



All-in-one
solution



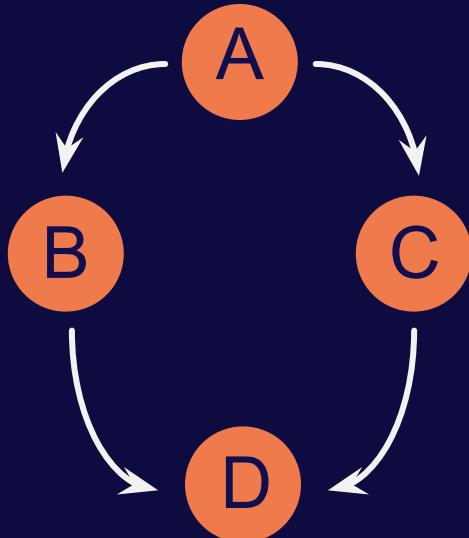
Extensive
documentation

- Business logic in functions
- Workflow orchestration logic of these functions
- Submission of workflows
- Walkthrough assumes zero knowledge
- User guides for more complex features

Enough talk, let's code!

```
pip install hera
```

From the Argo YAML experience...



```

apiVersion: argoproj.io/v1alpha1
kind: Workflow
metadata:
  generateName: dag-diamond-
spec:
  entrypoint: diamond
  templates:
  - name: diamond
    dag:
      tasks:
      - name: A
        template: echo
        arguments:
          parameters: [{name: message, value: A}]
    - depends: A
      name: B
      template: echo
      arguments:
        parameters: [{name: message, value: B}]
    - depends: A
      name: C
      template: echo
      arguments:
        parameters: [{name: message, value: C}]
    - depends: B && C
      name: D
      template: echo
      arguments:
        parameters: [{name: message, value: D}]
    - name: echo
      container:
        command: ["echo", "{{inputs.parameters.message}}"]
        image: alpine:3.7
      inputs:
        parameters:
        - name: message
  
```

...to Hera!

Hera provides many custom classes to help you author your Workflow - with code completion!

```
with Workflow(  
    ge  
):  
    [?] generate_name=  
    [?] generation=  
    [?] getattr  
    [?] GeneratorExit  
    [?] deletion_grace_period_seconds=
```

```
from hera.workflows import DAG, Container, Parameter, Task, Workflow  
  
with Workflow(generate_name="dag-diamond-", entrypoint="diamond") as w:  
    echo = Container(  
        name="echo",  
        image="alpine:3.7",  
        command=["echo", "{{inputs.parameters.message}}"],  
        inputs=[Parameter(name="message")],  
    )  
  
    with DAG(name="diamond"):  
        A = Task(name="A", template=echo, arguments={"message": "A"})  
        B = Task(name="B", template=echo, arguments={"message": "B"})  
        C = Task(name="C", template=echo, arguments={"message": "C"})  
        D = Task(name="D", template=echo, arguments={"message": "D"})  
        A >> [B, C] >> D
```

...to Hera!

Hera uses a context manager pattern for Workflows, DAGs, Steps and more, which mirrors the YAML syntax

```
from hera.workflows import DAG, Container, Parameter, Task, Workflow

with Workflow(generate_name="dag-diamond-", entrypoint="diamond") as w:
    echo = Container(
        name="echo",
        image="alpine:3.7",
        command=["echo", "{{inputs.parameters.message}}"],
        inputs=[Parameter(name="message")],
    )

    with DAG(name="diamond"):
        A = Task(name="A", template=echo, arguments={"message": "A"})
        B = Task(name="B", template=echo, arguments={"message": "B"})
        C = Task(name="C", template=echo, arguments={"message": "C"})
        D = Task(name="D", template=echo, arguments={"message": "D"})
        A >> [B, C] >> D
```

...to Hera!

Objects like this Container template are automatically detected and added to the Workflow's templates

```
from hera.workflows import DAG, Container, Parameter, Task, Workflow

with Workflow(generate_name="dag-diamond-", entrypoint="diamond") as w:
    echo = Container(
        name="echo",
        image="alpine:3.7",
        command=["echo", "{{inputs.parameters.message}}"],
        inputs=[Parameter(name="message")],
    )
    with DAG(name="diamond"):
        A = Task(name="A", template=echo, arguments={"message": "A"})
        B = Task(name="B", template=echo, arguments={"message": "B"})
        C = Task(name="C", template=echo, arguments={"message": "C"})
        D = Task(name="D", template=echo, arguments={"message": "D"})
        A >> [B, C] >> D
```

...to Hera!

```
from hera.workflows import DAG, Container, Parameter, Task, Workflow

with Workflow(generate_name="dag-diamond-", entrypoint="diamond") as w:
    echo = Container(
        name="echo",
        image="alpine:3.7",
        command=["echo", "{{inputs.parameters.message}}"],
        inputs=[Parameter(name="message")],
    )
    with DAG(name="diamond"):
        A = Task(name="A", template=echo, arguments={"message": "A"})
        B = Task(name="B", template=echo, arguments={"message": "B"})
        C = Task(name="C", template=echo, arguments={"message": "C"})
        D = Task(name="D", template=echo, arguments={"message": "D"})
        A >> [B, C] >> D
```

While these Tasks are added
to the DAG

With a sprinkling of syntactic sugar

```
from hera.workflows import DAG, Container, Parameter, Task, Workflow

with Workflow(generate_name="dag-diamond-", entrypoint="diamond") as w:
    echo = Container(
        name="echo",
        image="alpine:3.7",
        command=["echo", "{{inputs.parameters.message}}"],
        inputs=[Parameter(name="message")],
    )
    with DAG(name="diamond"):
        A = echo(name="A", arguments={"message": "A"})
        B = echo(name="B", arguments={"message": "B"})
        C = echo(name="C", arguments={"message": "C"})
        D = echo(name="D", arguments={"message": "D"})
        A >> [B, C] >> D
```

We can simply “call” the echo container template!

You can easily describe task dependencies with the `rshift` operator!

Container templates? I want more Python!

Functions are templates!

Use your functions as script templates with the script decorator!

```
from hera.workflows import DAG, Workflow, script

@script(image="python:3.12")
def echo(message):
    print(message)

with Workflow(generate_name="dag-diamond-", entrypoint="diamond") as w:
    with DAG(name="diamond"):
        A = echo(name="A", arguments={"message": "A"})
        B = echo(name="B", arguments={"message": "B"})
        C = echo(name="C", arguments={"message": "C"})
        D = echo(name="D", arguments={"message": "D"})
        A >> [B, C] >> D
```

Call them in your DAGs and Steps!

Create your workflow straight from your favourite IDE!

```
from hera.workflows import DAG, Workflow, WorkflowsService\n\nwith Workflow(\n    generate_name="dag-diamond-",\n    entrypoint="diamond",\n    namespace="argo",\n    workflows_service=WorkflowsService(\n        host="https://localhost:2746"),\n) as w:\n    with DAG(name="diamond"):\n        ...\n    w.create()
```

Add some info to communicate with your Argo instance – your namespace and a **WorkflowsService**!

Call **create** on your workflow!



A script template caveat...

Hera “compiles” your workflow definition to YAML for Argo to understand it

Script templates expect the code dumped into the **source** field of the YAML

In Hera, we call this an “inline” script because the code is “inline” in your YAML

```
from hera.workflows import script

@script(image="python:3.12")
def echo(message):
    print(message)

- name: echo
  inputs:
    parameters:
      - name: message
  script:
    command:
      - python
    image: python:3.12
    source: |-
      import json
      try: message = json.loads(r'''{{inputs.parameters.message}}''')
      except: message = r'''{{inputs.parameters.message}}'''

    print(message)
```

A script template caveat...

So, our function is **type-less** and practically **untestable**!

```
from hera.workflows import script

@script(image="python:3.12")
def echo(message):
    print(message) # returns are not allowed!

- name: echo
  inputs:
    parameters:
      - name: message
  script:
    command:
      - python
    image: python:3.12
    source: |-
      import json
      try: message = json.loads(r'''{{inputs.parameters.message}}''')
      except: message = r'''{{inputs.parameters.message}}'''

    print(message)
```

Hera helps by adding **json.loads** for your parameters, but we're at the mercy of whatever the user gives us!



What if we could
do better?

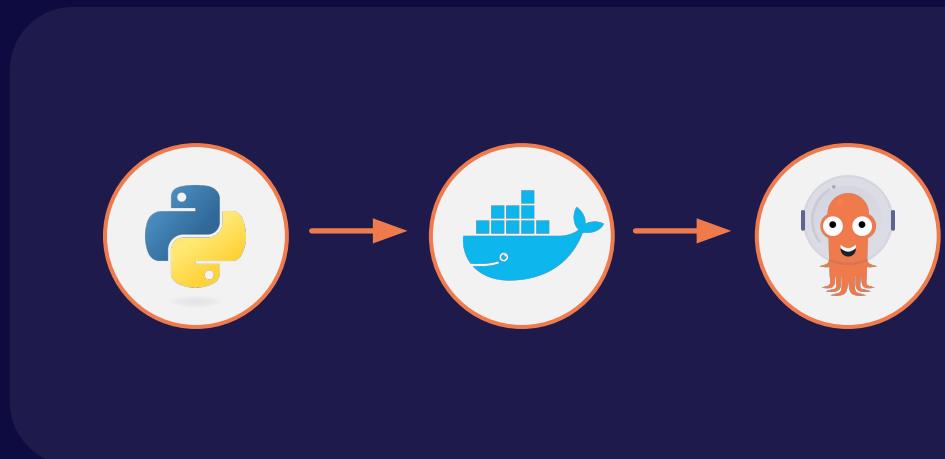
The Hera Script Runner

Hera's Script Runner runs on an image built using your code

- ✓ Users can build a Buildpacks/Docker image with their code and dependencies
- ✓ Argo then pulls that image to run the code as written

The YAML that Hera exports uses some magic to make this happen!

The command is still Python, but the source is now the input parameters which are passed to the `hera.workflows.runner` module



The Hera Script Runner

Hera's Script Runner runs on an image built using your code

- ✓ Users can build a Buildpacks/Docker image with their code and dependencies
- ✓ Argo then pulls that image to run the code as written

The YAML that Hera exports uses some magic to make this happen!

The command is still Python, but the source is now the input parameters which are passed to the `hera.workflows.runner` module



```
script:  
  image: my-built-python-image  
  command:  
    - python  
  args:  
    - -m  
    - hera.workflows.runner  
    - -e  
    - argocon_tutorial:calculate_area_of_rectangle  
  source: '{{inputs.parameters}}'
```

Type-safe functions!

We first need to specify the **constructor** as “runner”

And we need our template to run on an **image** that we’re going to build from this code

```
from hera.workflows import script

@script(
    constructor="runner",
    image="my-built-python-image",
)
def calculate_area_of_rectangle(
    length: float, width: float
) -> float:
    return length * width
```

Type-safe functions!

Parameters are still inferred from the function signature and exported to YAML

The Hera runner deserializes and type checks the inputs at runtime!

```
- name: calculate-area-of-rectangle
  inputs:
    parameters:
      - name: length
      - name: width
    command:
      - python
  script:
    args:
      - -m
      - hera.workflows.runner
      - -e
      - argocon_tutorial:calculate_area_of_rectangle
    image: my-built-python-image
    source: '{{inputs.parameters}}'
```

```
from hera.workflows import script

@script(
    constructor="runner",
    image="my-built-python-image",
)
def calculate_area_of_rectangle(
    length: float, width: float
) -> float:
    return length * width
```



Type-safe functions!

With the Hera runner, we can even return values straight from the function!

Which means...

```
from hera.workflows import script

@script(
    constructor="runner",
    image="my-built-python-image",
)
def calculate_area_of_rectangle(
    length: float, width: float
) -> float:
    return length * width
```

Type-safe functions – testing!

```
from hera.workflows import script

@script(
    constructor="runner",
    image="my-built-python-image",
)
def calculate_area_of_rectangle(
    length: float, width: float
) -> float:
    return length * width
```

We can test script template functions
with normal Python tests!

```
def test_calculate_area_of_rectangle():
    assert calculate_area_of_rectangle(2.0, 3.0) == 6.0
```

Type-safe functions – testing!

However, only JSON types are allowed in the inputs...

```
from hera.workflows import script

@script(
    constructor="runner",
    image="my-built-python-image",
)
def calculate_area_of_rectangle(
    length: float, width: float
) -> float:
    return length * width

def test_calculate_area_of_rectangle():
    assert calculate_area_of_rectangle(2.0, 3.0) == 6.0
```

What if we could do...
even better?

Hera ❤️ Pydantic



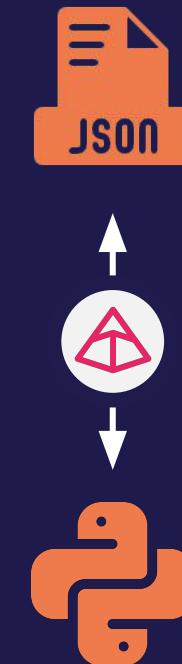
```
pip install pydantic
```

Pydantic-powered functions!

The Hera Script Runner integrates with Pydantic for type validation, and allows Pydantic classes in your template inputs and outputs

- You get Python type hints and auto-completion during template development
- The Script Runner gets your Pydantic objects for you by automatically deserializing JSON string inputs

Ensuring type safety for input objects at runtime!



Pydantic-powered functions!

Let's create a Pydantic
BaseModel class

And we can give it a function

```
from pydantic import BaseModel

class Rectangle(BaseModel):
    length: float
    width: float

    def area(self) -> float:
        return self.length * self.width
```

Pydantic-powered functions!

Pass Pydantic types directly into
(and out of) functions!

Unit test your template logic just
like any other function!

```
class Rectangle(BaseModel):
    length: float
    width: float

    def area(self) -> float:
        return self.length * self.width
```

```
@script(constructor="runner", image="my-built-python-image")
def calculate_area_of_rectangle(rectangle: Rectangle) -> float:
    return rectangle.area()
```

```
def test_calculate_area_of_rectangle():
    r = Rectangle(length=2.0, width=3.0)
    assert calculate_area_of_rectangle(r) == r.area()
```

I like testing things!

Good news, everyone! Workflows are testable!

We can launch a workflow,
waiting until completion...

```
with Workflow(generate_name="hello-world-", entrypoint="steps") as w:
    with Steps(name="steps"):
        echo_to_param(arguments={"message": "Hello world!"})

    def test_create_workflow():
        model_workflow = w.create(wait=True)
        assert model_workflow.status and
            model_workflow.status.phase == "Succeeded"

        echo_node = next(filter(
            lambda n: n.display_name == "echo-to-param",
            model_workflow.status.nodes.values(),
        ))

        assert echo_node.outputs.parameters[0].value == "Hello world!"
```

Good news, everyone!

Workflows are testable!

Inspect the Workflow's status...

```
with Workflow(generate_name="hello-world-", entrypoint="steps") as w:
    with Steps(name="steps"):
        echo_to_param(arguments={"message": "Hello world!"})

    def test_create_workflow():
        model_workflow = w.create(wait=True)
        assert model_workflow.status and
            model_workflow.status.phase == "Succeeded"

        echo_node = next(filter(
            lambda n: n.display_name == "echo-to-param",
            model_workflow.status.nodes.values(),
        ))

        assert echo_node.outputs.parameters[0].value == "Hello world!"
```

Good news, everyone! Workflows are testable!

```
with Workflow(generate_name="hello-world-", entrypoint="steps") as w:
    with Steps(name="steps"):
        echo_to_param(arguments={"message": "Hello world!"})

    def test_create_workflow():
        model_workflow = w.create(wait=True)
        assert model_workflow.status and
            model_workflow.status.phase == "Succeeded"

        echo_node = next(filter(
            lambda n: n.display_name == "echo-to-param",
            model_workflow.status.nodes.values(),
        ))
        assert echo_node.outputs.parameters[0].value == "Hello world!"
```

Get individual nodes...

And check their outputs!

Hera's future goals

- Flag icon Bring automatic Python versioning to WorkflowTemplates
- Flag icon Leverage Python dependency management to manage WorkflowTemplate dependencies



How is Hera used?

Bloomberg

Problem

Building and testing complex ML model remediation pipelines



- Argo Workflows has all the features we need for orchestrating these pipelines
- It is provided as a service at Bloomberg, but was not adopted for ML use cases
 - Difficult to write and test Workflows with YAML
 - YAML workflows had Python code dumped into script templates

Solution

Provide a Python experience for Argo Workflows

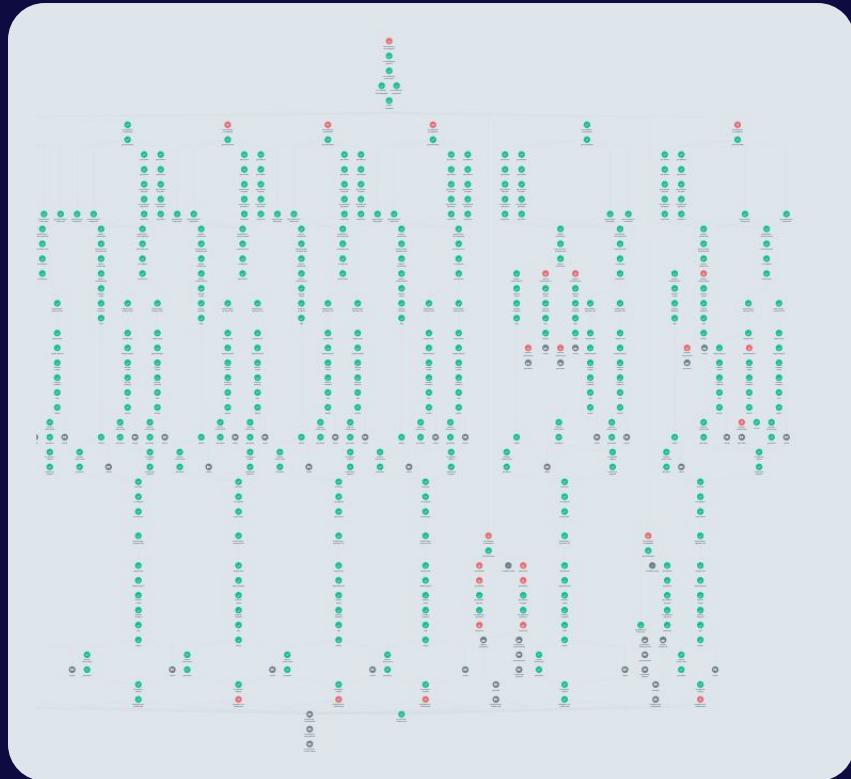


- Work on and release Hera V5
- Provide supporting tools
 - Cookiecutter to provide a boilerplate Python package
 - Testing framework
 - Documentation generator for Workflow Templates
- Workshops and support channels

TechAtBloomberg.com

Bloomberg

Hera has enabled our 300+
AI engineers to use Argo
Workflows in production



TechAtBloomberg.com

© 2024 Bloomberg Finance L.P. All rights reserved.

Bloomberg
Engineering

Pipekit

PROBLEM

managing Argo at scale

- Monitoring/observability for Workflows
- Manage Argo/Hera at scale and across multiple clusters

SOLUTION

tools and support for managing Argo at scale

- Providing easy integrations with common tools
- **pipekit-sdk** - builds on Hera to allow users to submit to multiple clusters, interact with workflows, and see logs from a single Jupyter Notebook or any Python interpreter

LEARNINGS

- Need to meet both the platform team and data scientists where they're at
- Platform teams need to use normal K8s primitives and declarations they're used to
- Data scientists want to use languages and concepts they're used to

ACCURE Battery Intelligence

Problem

YAML and slow Docker builds
blocking productivity



Solution

Adopt Hera to improve
developer experience

- Difficult to set resources and limits for Python developers
- Unaware of what parameters they can play with
- Had to rebuild the Docker image every time code was changed

- Hera helps with discoverability of Workflows parameters
- Using Hera's "inline" scripts allows developers to avoid rebuilding the image

Energy Company

Problem

Markets moving too fast for slow iteration speeds



- Data scientists are familiar with Python and not YAML
- Required help from MLOps team to run workflows
 - MLOps team blocked from details by RBAC policies

Solution

Adopt Hera to allow the data science team to be self serve



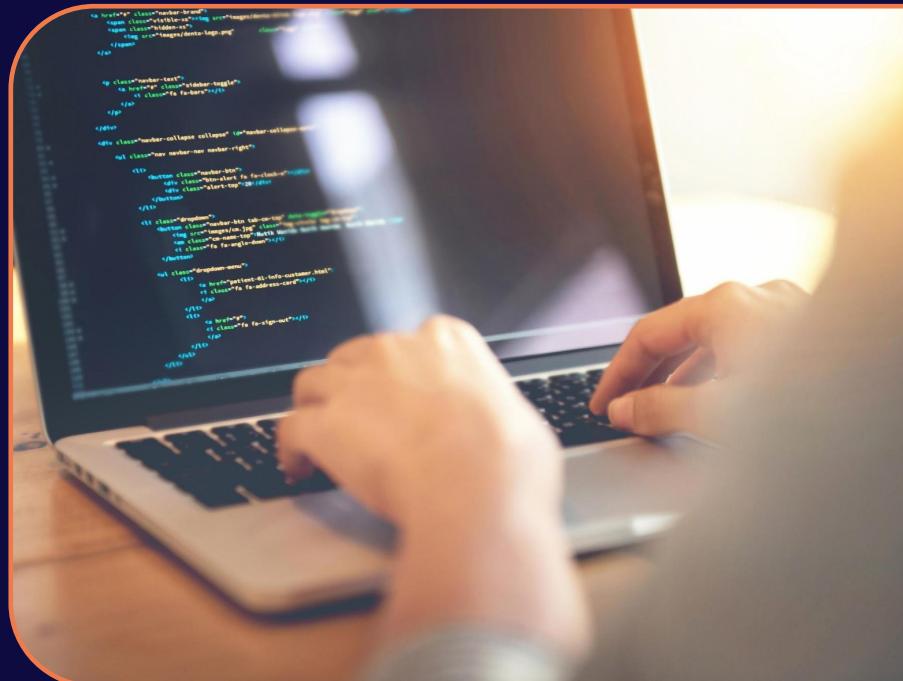
- All-in-one solution for data scientists to write training code, orchestration logic, and submission to the Argo Workflows clusters

Key Takeaways

- ✓ Hera makes Argo Workflows easy
- ✓ Argo does the heavy lifting of orchestrating your Python functions
- ✓ Hera is used in production, operating at scale

TL;DR:

Argo Workflows and Hera can give you the best cloud-native workflow orchestration experience



Thank you & Questions



Repository

github.com/argoproj-labs/hera

Docs

hera.readthedocs.io



#hera-argo-sdk channel

on slack.cncf.io

*Working Group meetings
coming soon!*