

OPEN
AI + DATA
FORUM



Automating Cloud-native Spark Jobs with Argo Workflows

Caelan Urquhart & Darko Janjić, Pipekit

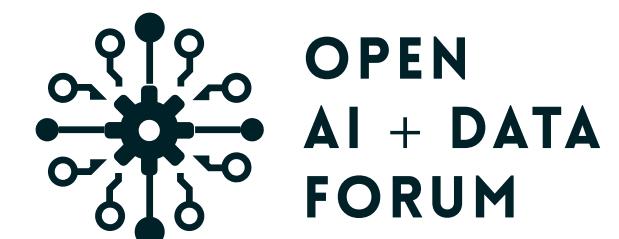
#OSSUMMIT

@pipekitio

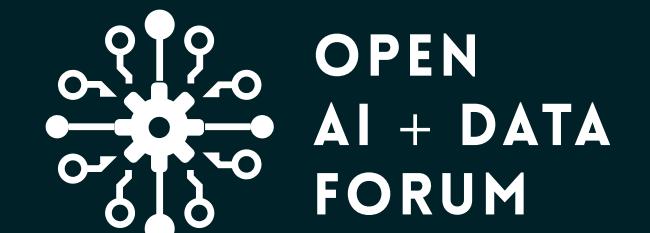


AGENDA

- **Intro & goals for the talk**
- **Background: Apache Spark, Kubernetes, and Argo Workflows**
- **Common problems with Spark**
- **Solutions with Kubernetes and Argo**
- **Demo & How-to**
- **Next steps**
- **Q&A**



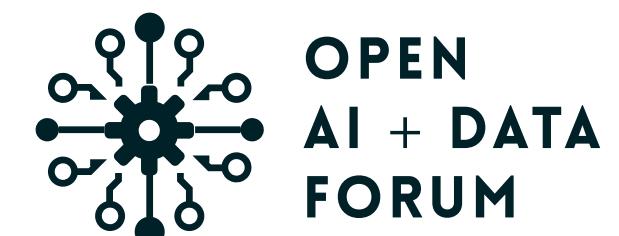
Intro & goals for the talk



#OSSUMMIT

GOALS FOR THE TALK

- Learn about the scaling and stability advantages of running Spark jobs on Kubernetes
- Learn how to use Argo Workflows to deploy and automate Spark jobs on Kubernetes
- Learn what else is possible once you're running Spark on Kubernetes



ABOUT US



Caelan Urquhart

Co-founder, CEO @ [Pipekit](#)



Contributer to the [Argo Workflows](#) project

Focus: Data engineering architecture,
product development



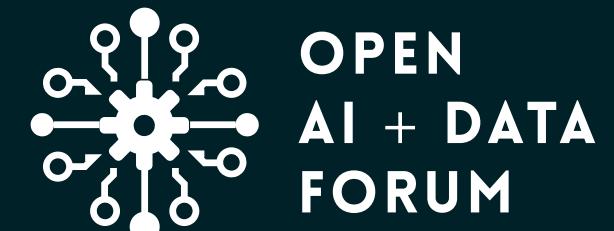
Darko Janjić

Senior Software Engineer @ [Pipekit](#)



Focus: Distributed systems, virtualization,
and cloud infrastructure engineering

Background: Apache Spark, Kubernetes, and Argo Workflows



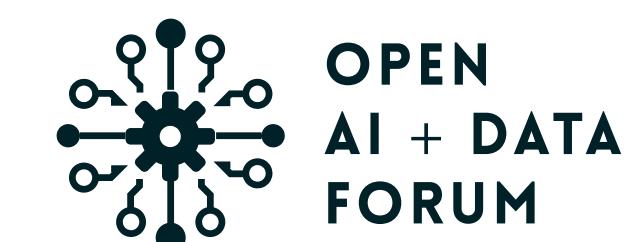
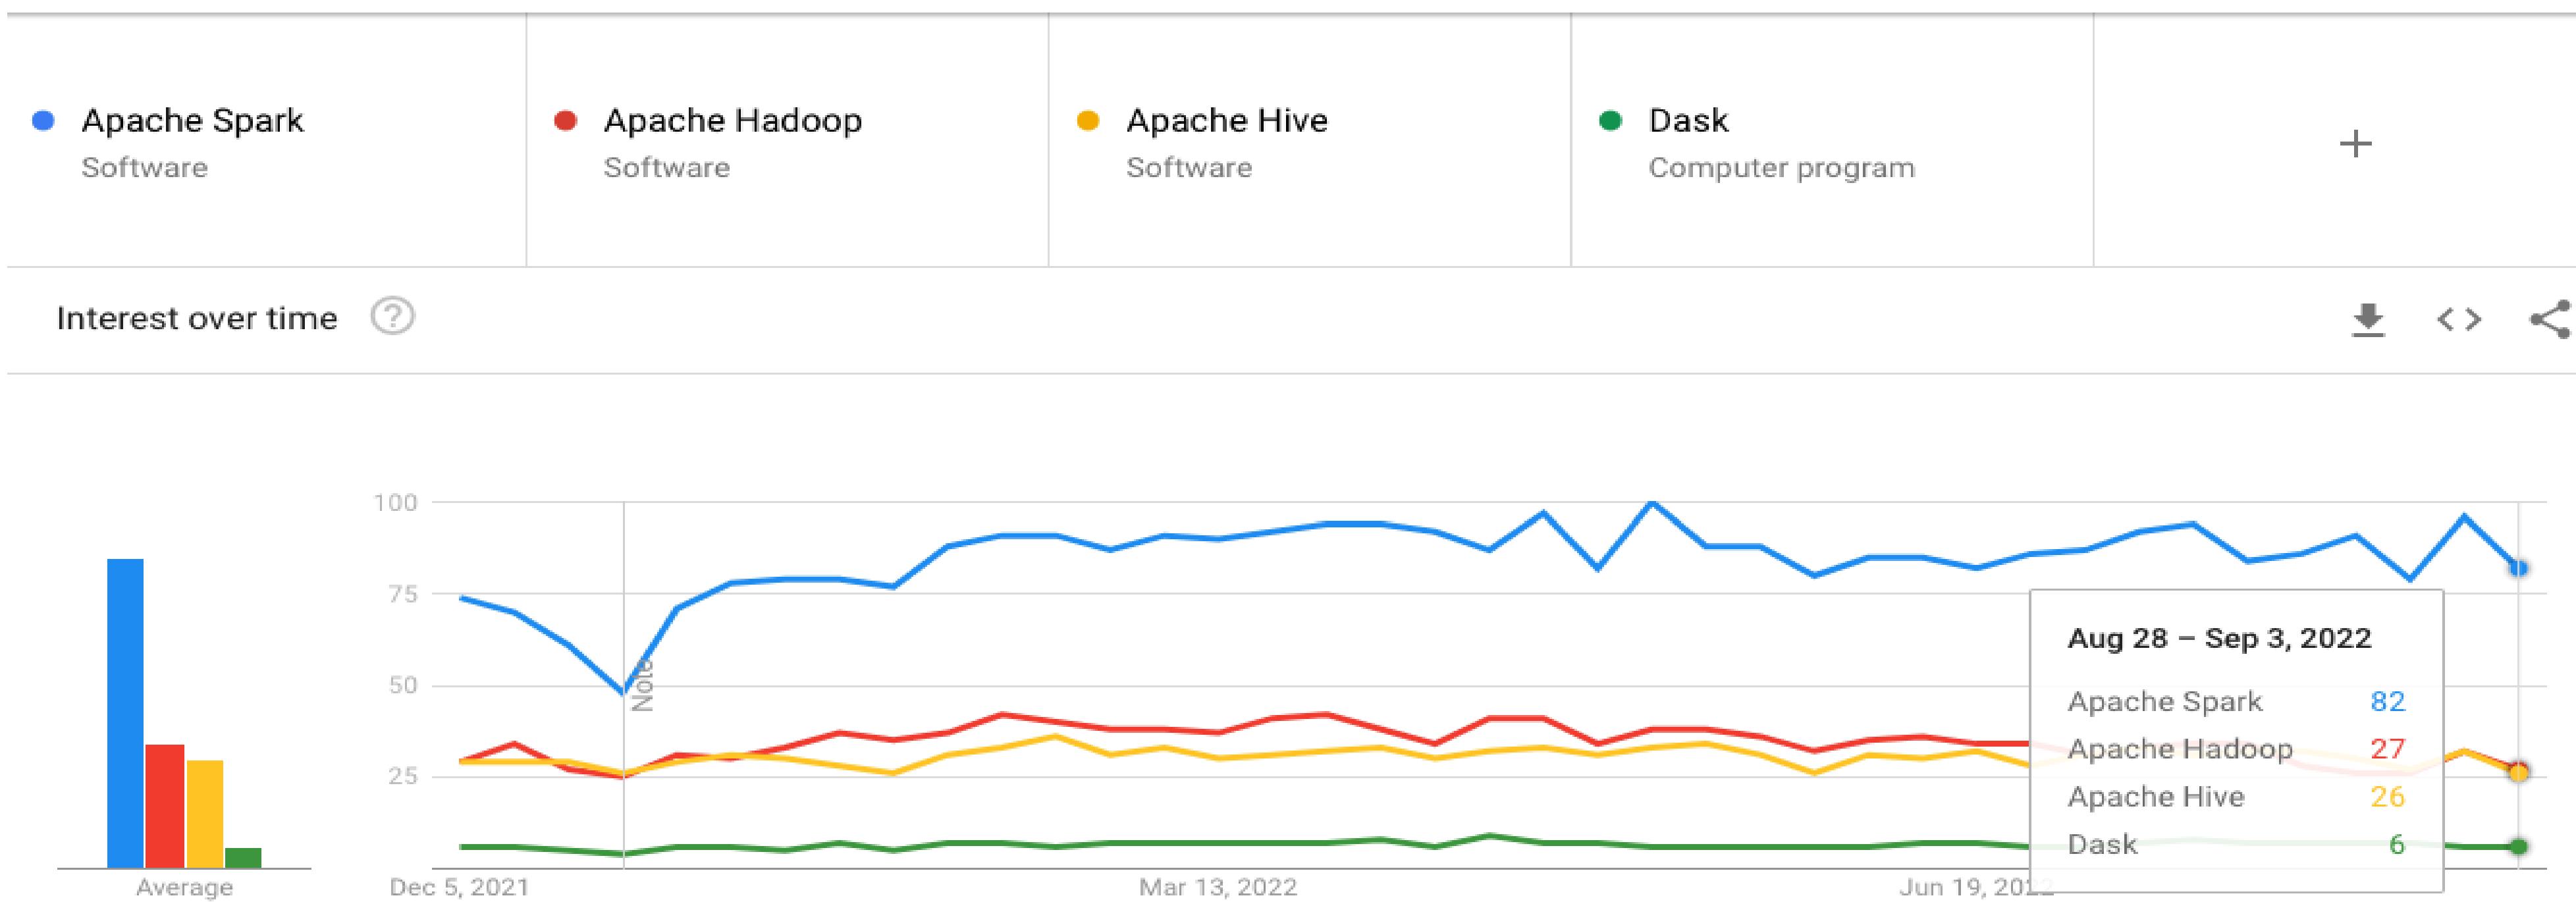
#OSSUMMIT

BACKGROUND: APACHE SPARK



The most popular big data compute framework

- Q: Why? A: Versatility
 - Structured Streaming, MLlib
 - Use Python, SQL, Java, Scala or R



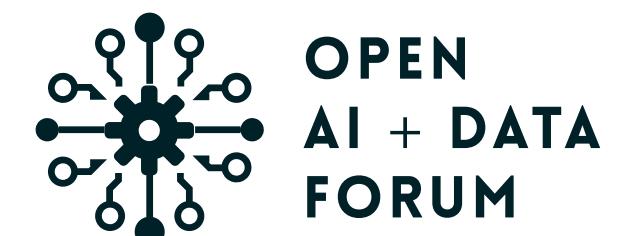
PROBLEMS

◆ The typical Spark deployment using YARN ...

- Requires global installs (cluster for Spark 3.0 vs. Spark 3.3)
- Doesn't support Docker natively
- Consumes more resources (JVM on each node)
- No autoscaling

◆ ... makes solving these problems frustrating

- How do I manage dependencies between different libraries or frameworks?
- How can I run experiments across local and prod environments?
- How can I use different GPUs on one cluster?
- How do I provision resources for “spikey” data jobs?

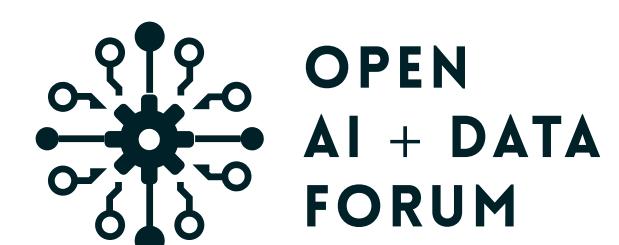


BACKGROUND: KUBERNETES

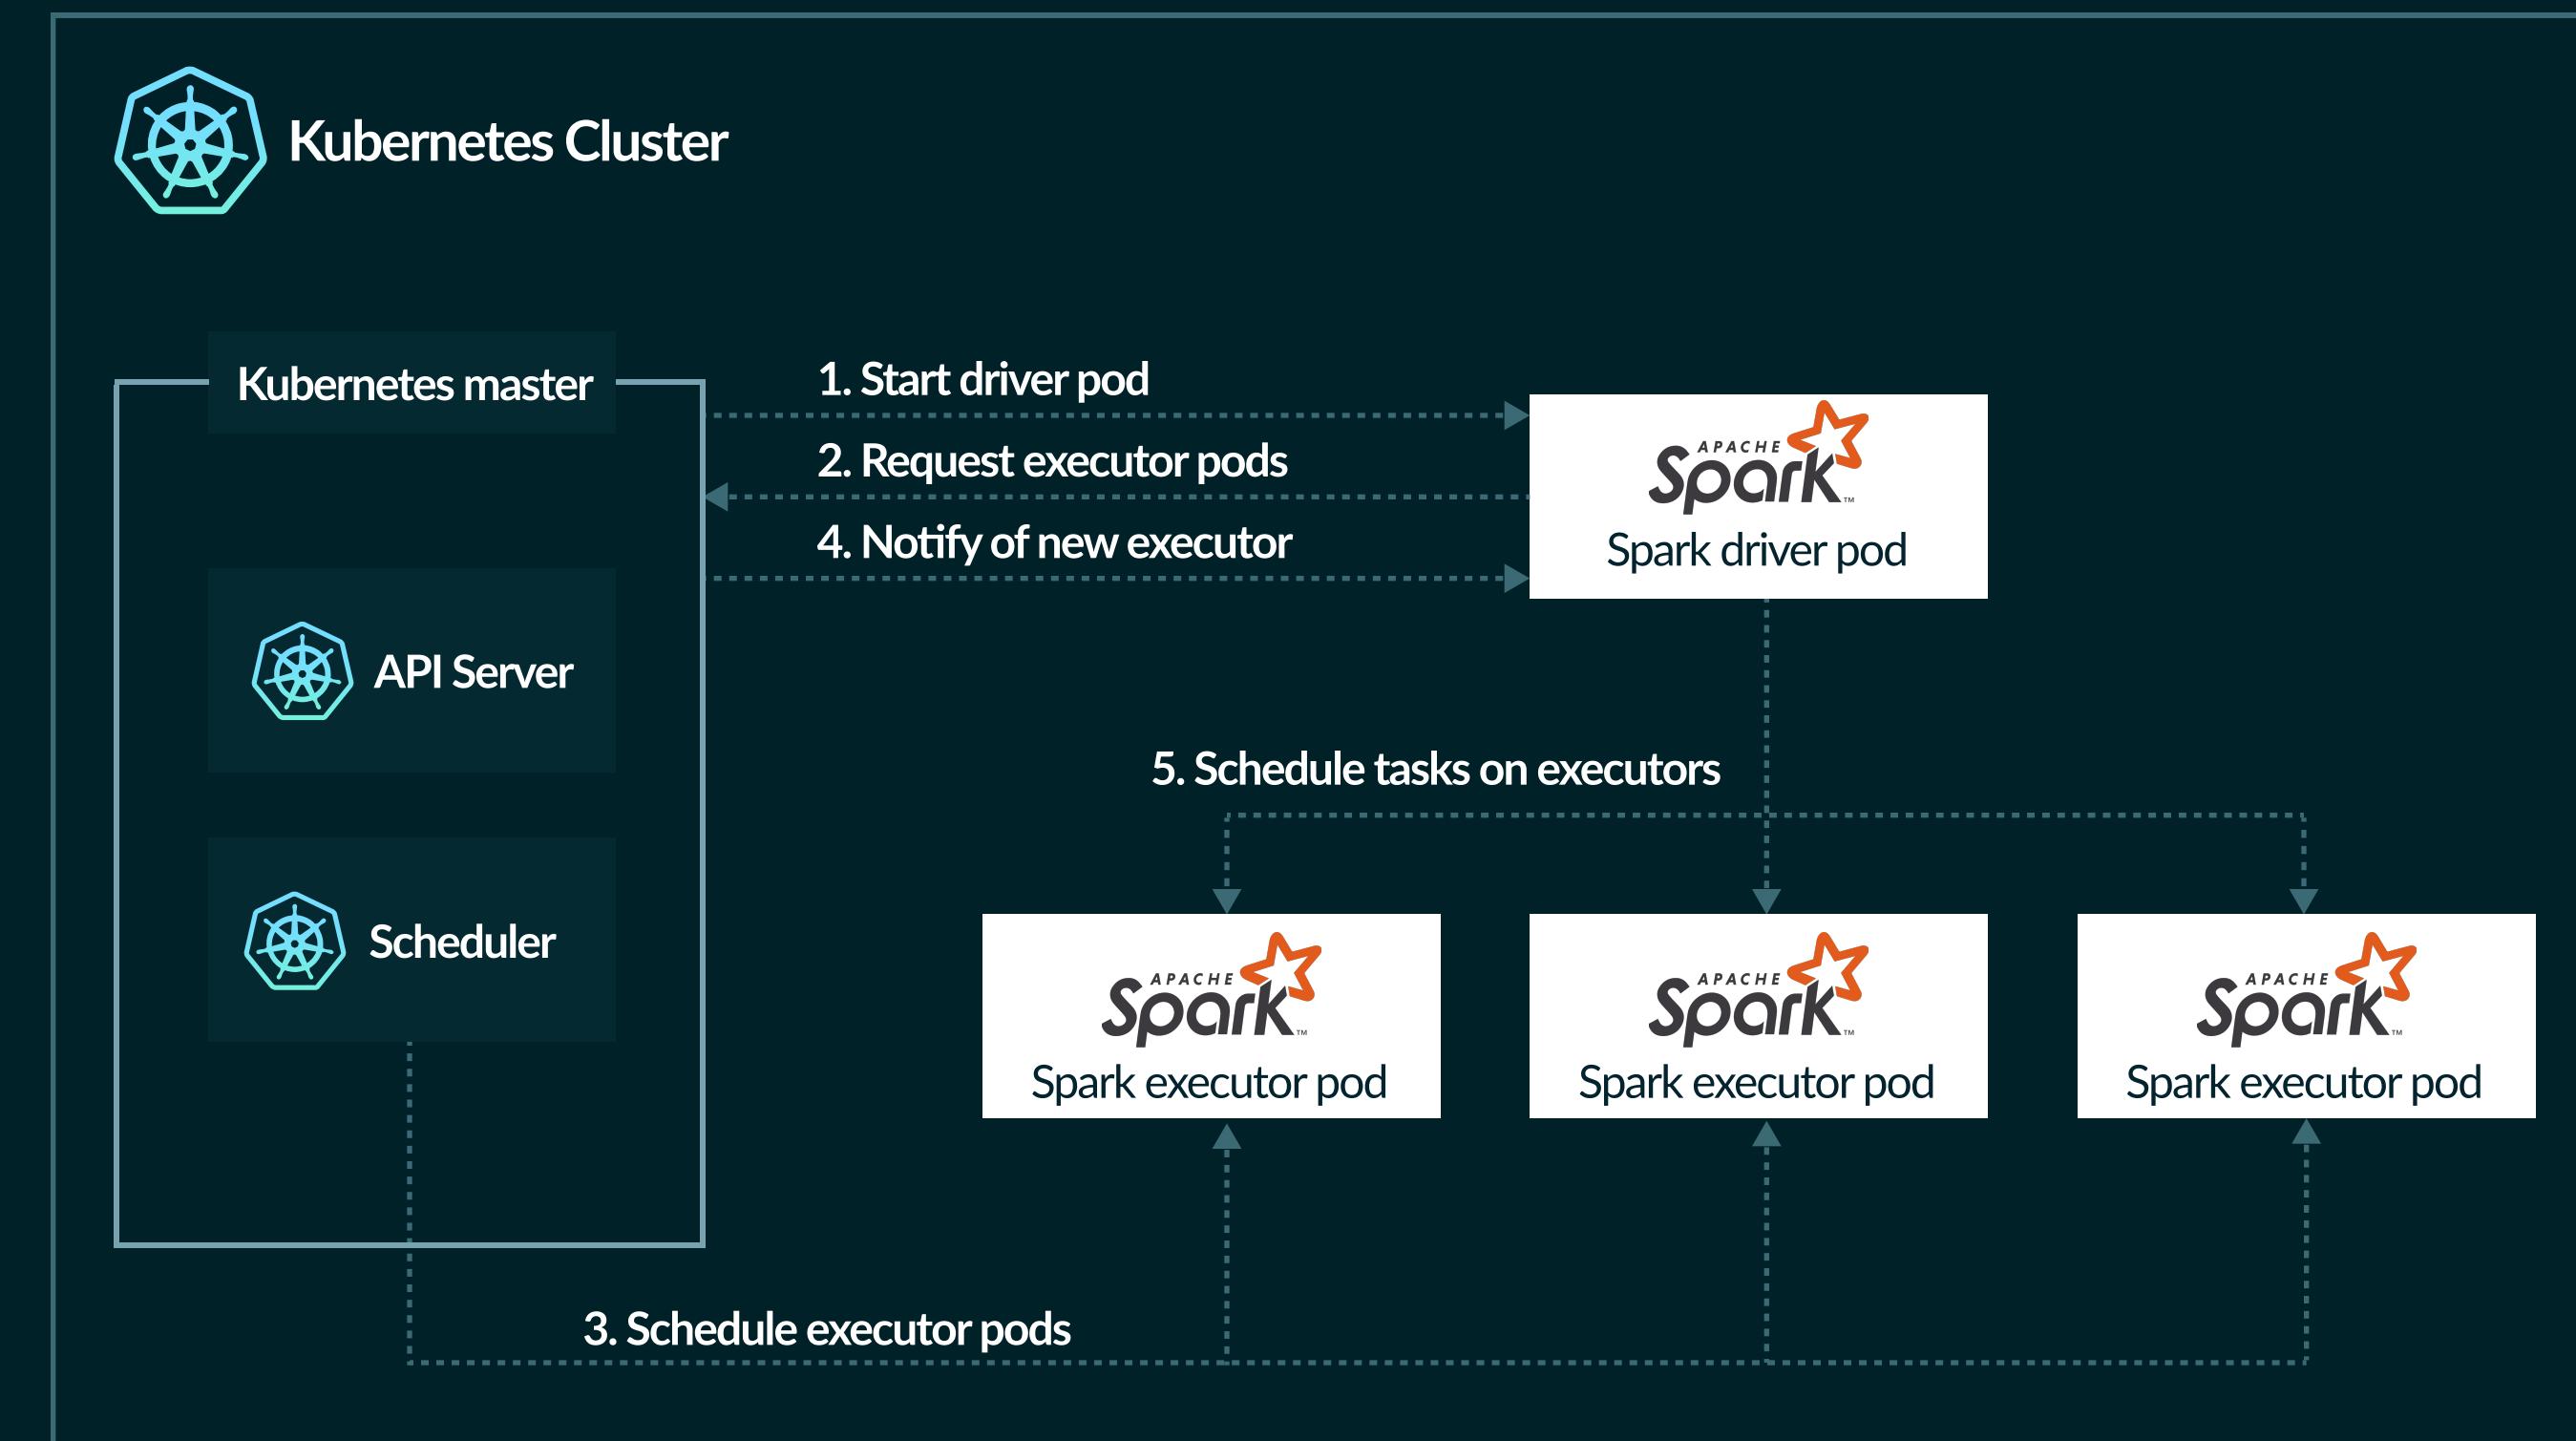
→ Kubernetes is a framework for managing and automating containerized applications.

Why use Kubernetes for data science?

- When **containers** are your building blocks you get...
 - Reproducibility
 - Reliability
- **Declarative** approach to data pipelines
- **Vertical autoscaling**: handle “spikey” data jobs
- Tap into the **cloud-native ecosystem** for additional tooling:
CI/CD, observability, etc.



SPARK ON K8S ARCHITECTURE



BACKGROUND: ARGO WORKFLOWS

Argo Workflows is the best way to run workloads on K8s

It's **generalizable** - use it to
automate anything:

- ETL, ELT, batch data processing
- ML training, serving
- CI/CD workloads

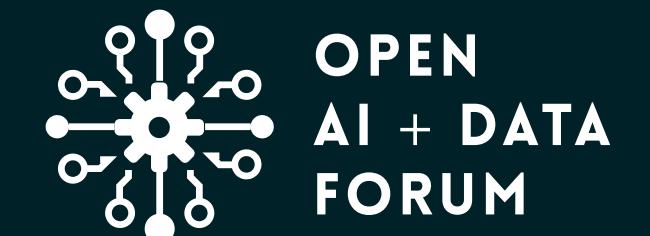
It's a workflow engine (aka
DAG orchestrator) that's
lightweight to deploy



Use **YAML or Python** to define
workflows

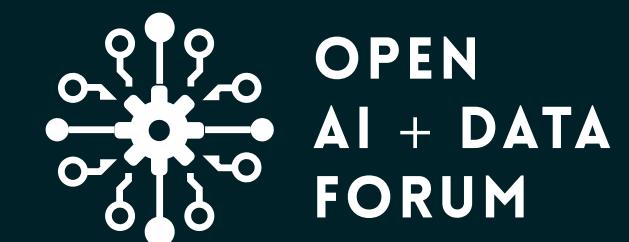
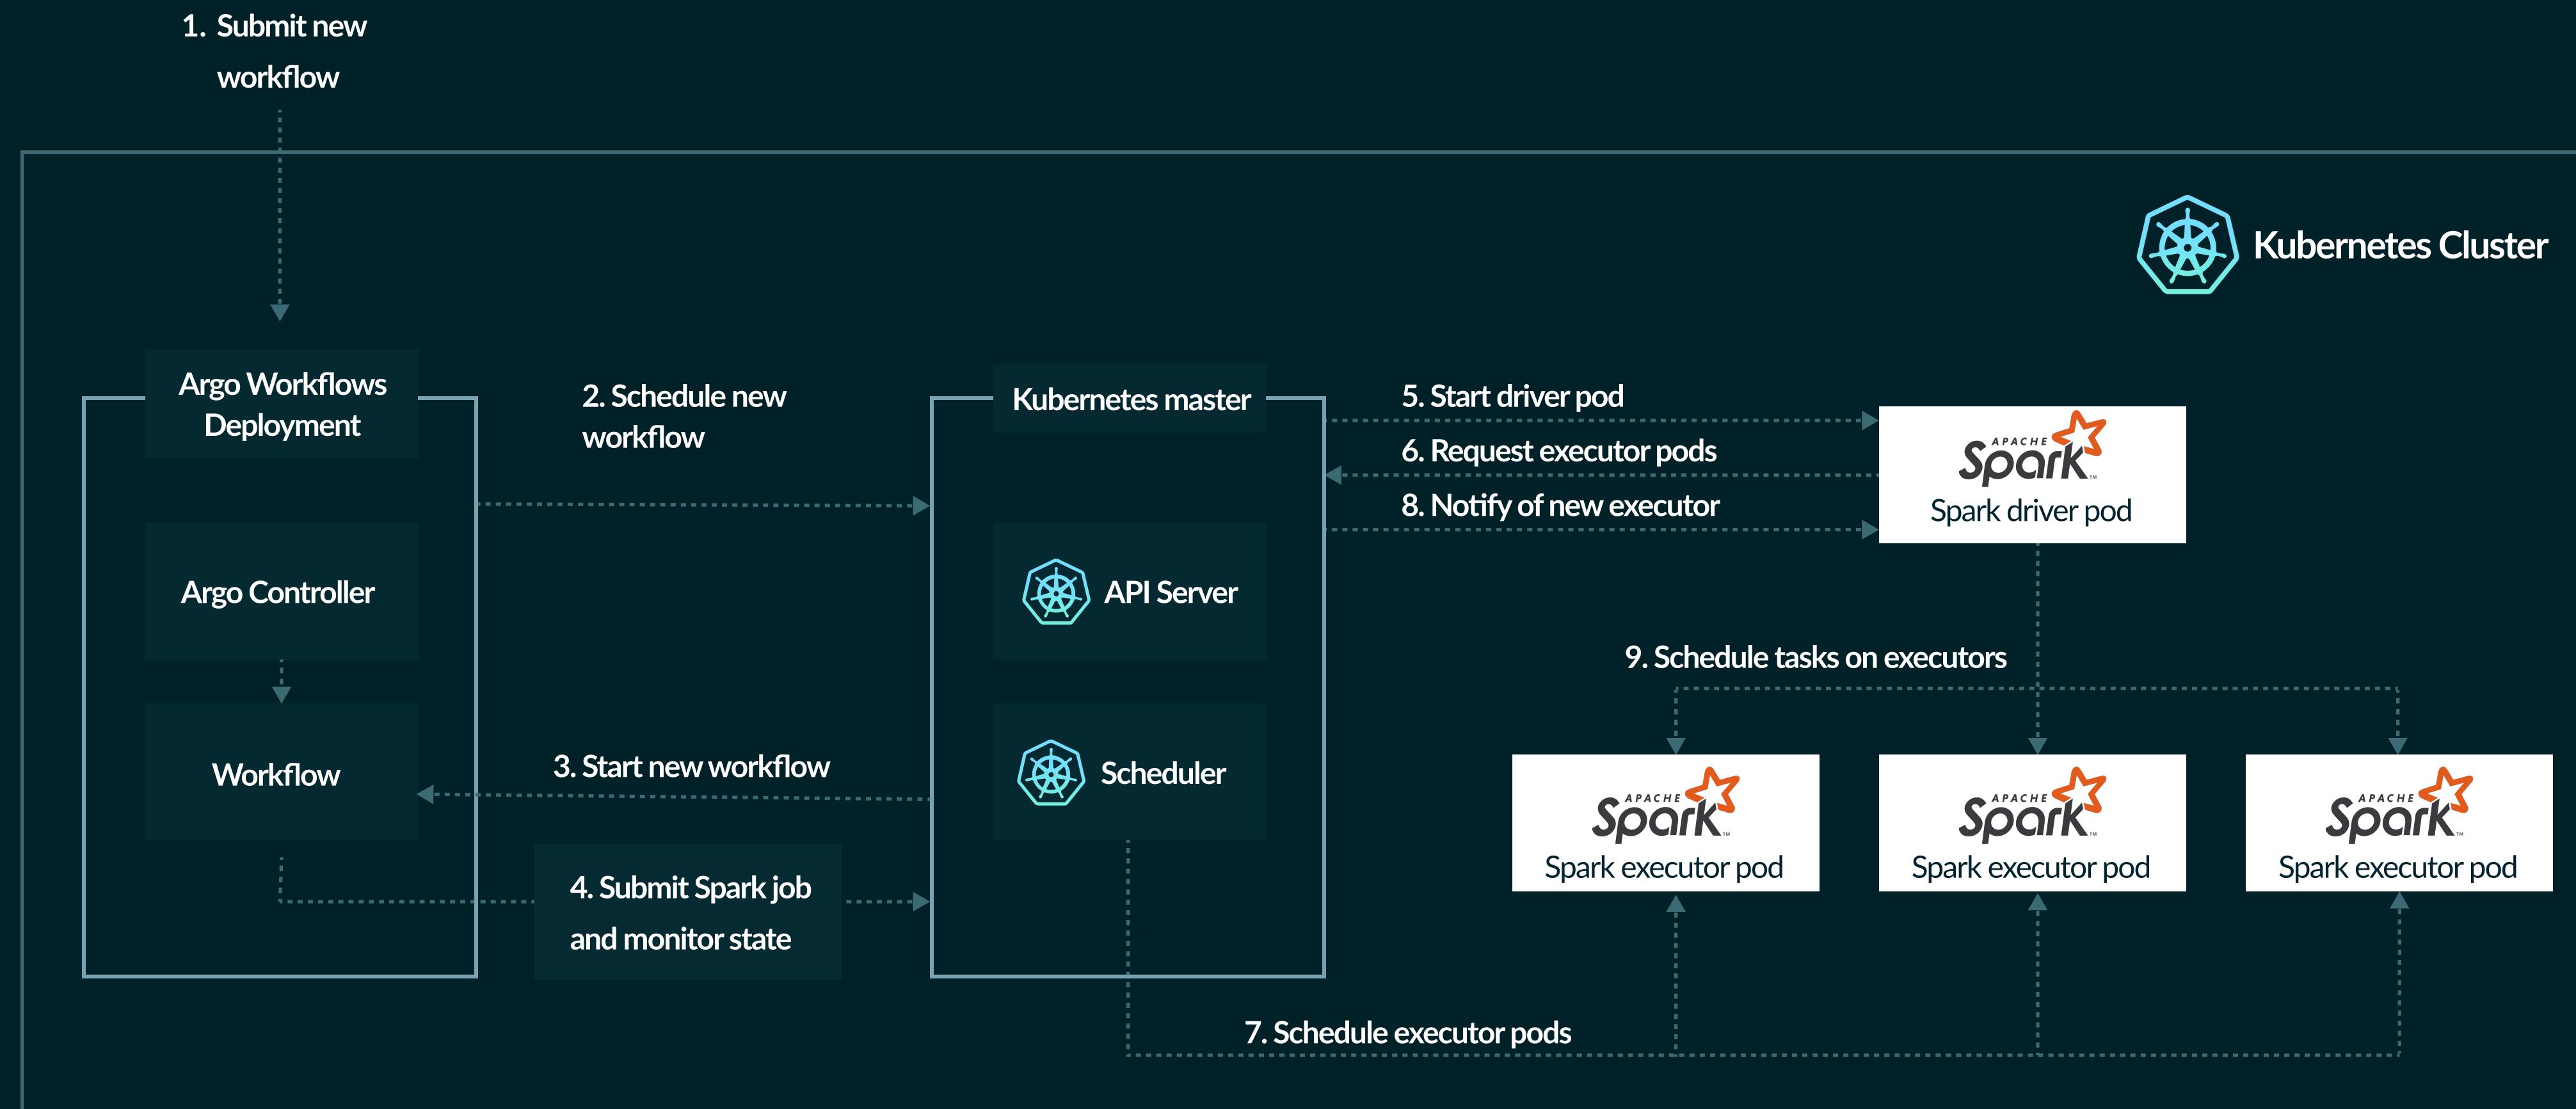
CNCF open source project

Solution



#OSSUMMIT

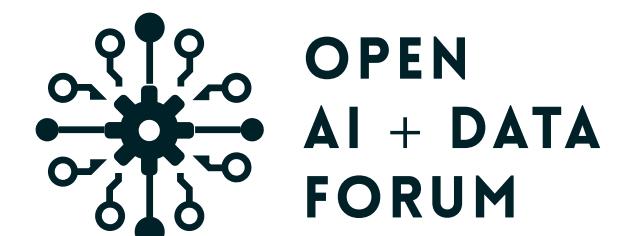
SPARK ON K8S ARCHITECTURE WITH ARGO



SOLUTION

→ Deploying Spark with K8s and Argo Workflows solves many of the issues with a traditional Spark deployment. ←

- **Container-native:**
 - run different Spark versions for different jobs on the same cluster
 - manage dependencies easily
 - stop downloading dependencies at runtime = faster devex
 - standardize environments to enable reproducibility
- **Lower resource requirements**
- **Autoscale your compute vertically**

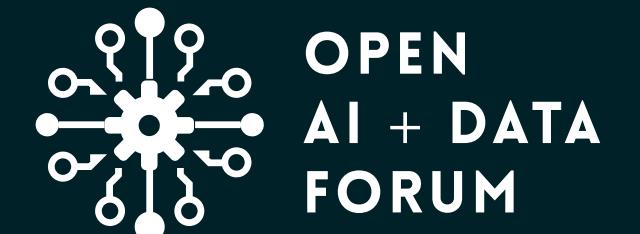


SOLUTION

➡ But wait, there's more! Adopting K8s and Argo Workflows opens up bonus benefits:

- **Duplicable**: use a similar architecture for other frameworks: Dask, Ray, etc.
- **Extendable**: bring other data processing and ML tooling into the same cluster: Apache Kafka, Beam, etc.
- **Cloud-native ecosystem**: add logging, metrics, and more
- **Scalable**: enable CI/CD, GitOps for your data platform
- **Agnostic**: avoid cloud vendor lock-in

Demo & How-to



#OSSUMMIT

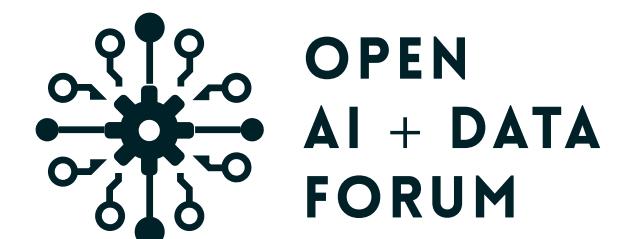
DEMO & HOW-TO

→ In this demo, we will...

- Create a Spark job that analyses a Kaggle dataset
- Create and run an Argo Workflow that executes the Spark job on K8s
- Simplify the Argo Workflow manifest by using a WorkflowTemplate

→ Prerequisites:

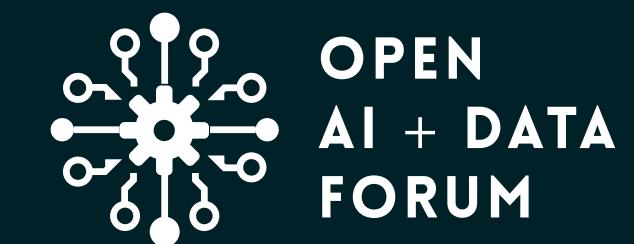
- K8s cluster running locally (K3d, Docker Desktop, Minikube, etc.)
- Spark operator installed
- Argo Workflows v3.3 installed
- Assumes you already have a Docker image for your Spark job



DEMO & HOW-TO

Spark job (Scala):

```
import org.apache.spark.sql.SparkSession  
  
import java.util.concurrent.TimeUnit  
  
object BikeRideLengthJob {  
  
    val fileLocation = "/app/df_1_year.csv"  
  
    def main(args: Array[String]): Unit = {  
        Logger.getLogger("org").setLevel(Level.ERROR)  
  
        val spark = SparkSession  
            .builder  
            .appName("BikeRideLengthJob")  
            .getOrCreate()  
  
        val rdd = spark.sparkContext.textFile(fileLocation).map(Bike.toBikeRide)  
  
        val bikeTypeStats = rdd  
            .map(x => (x.bikeType, x.rideLength))  
            .reduceByKey((x, y) => x + y)  
            .sortBy(x => x._2, ascending = false).collect()  
  
        for (elem <- bikeTypeStats) {  
            println(s"BikeType ${elem._1}, rideLength in hours ${TimeUnit.MILLISECONDS.toHours(elem._2)}")  
        }  
  
        spark.stop()  
    }  
}
```



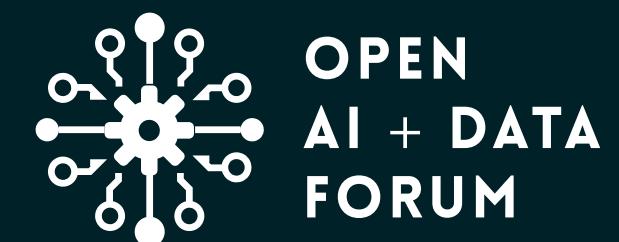
DEMO & HOW-TO

Argo Workflows manifest for our Spark jobs as a DAG:

```
apiVersion: argoproj.io/v1alpha1
kind: Workflow
metadata:
  generateName: scala-argo-spark-
  namespace: argo
spec:
  entrypoint: dag
  templates:
    - name: dag
      dag:
        tasks:
          - name: bike-type
            template: bike-type
          - name: bike-ride-length
            template: bike-ride-length
```

Defining the resources for our Spark job:

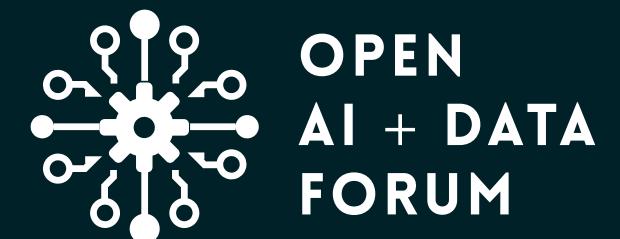
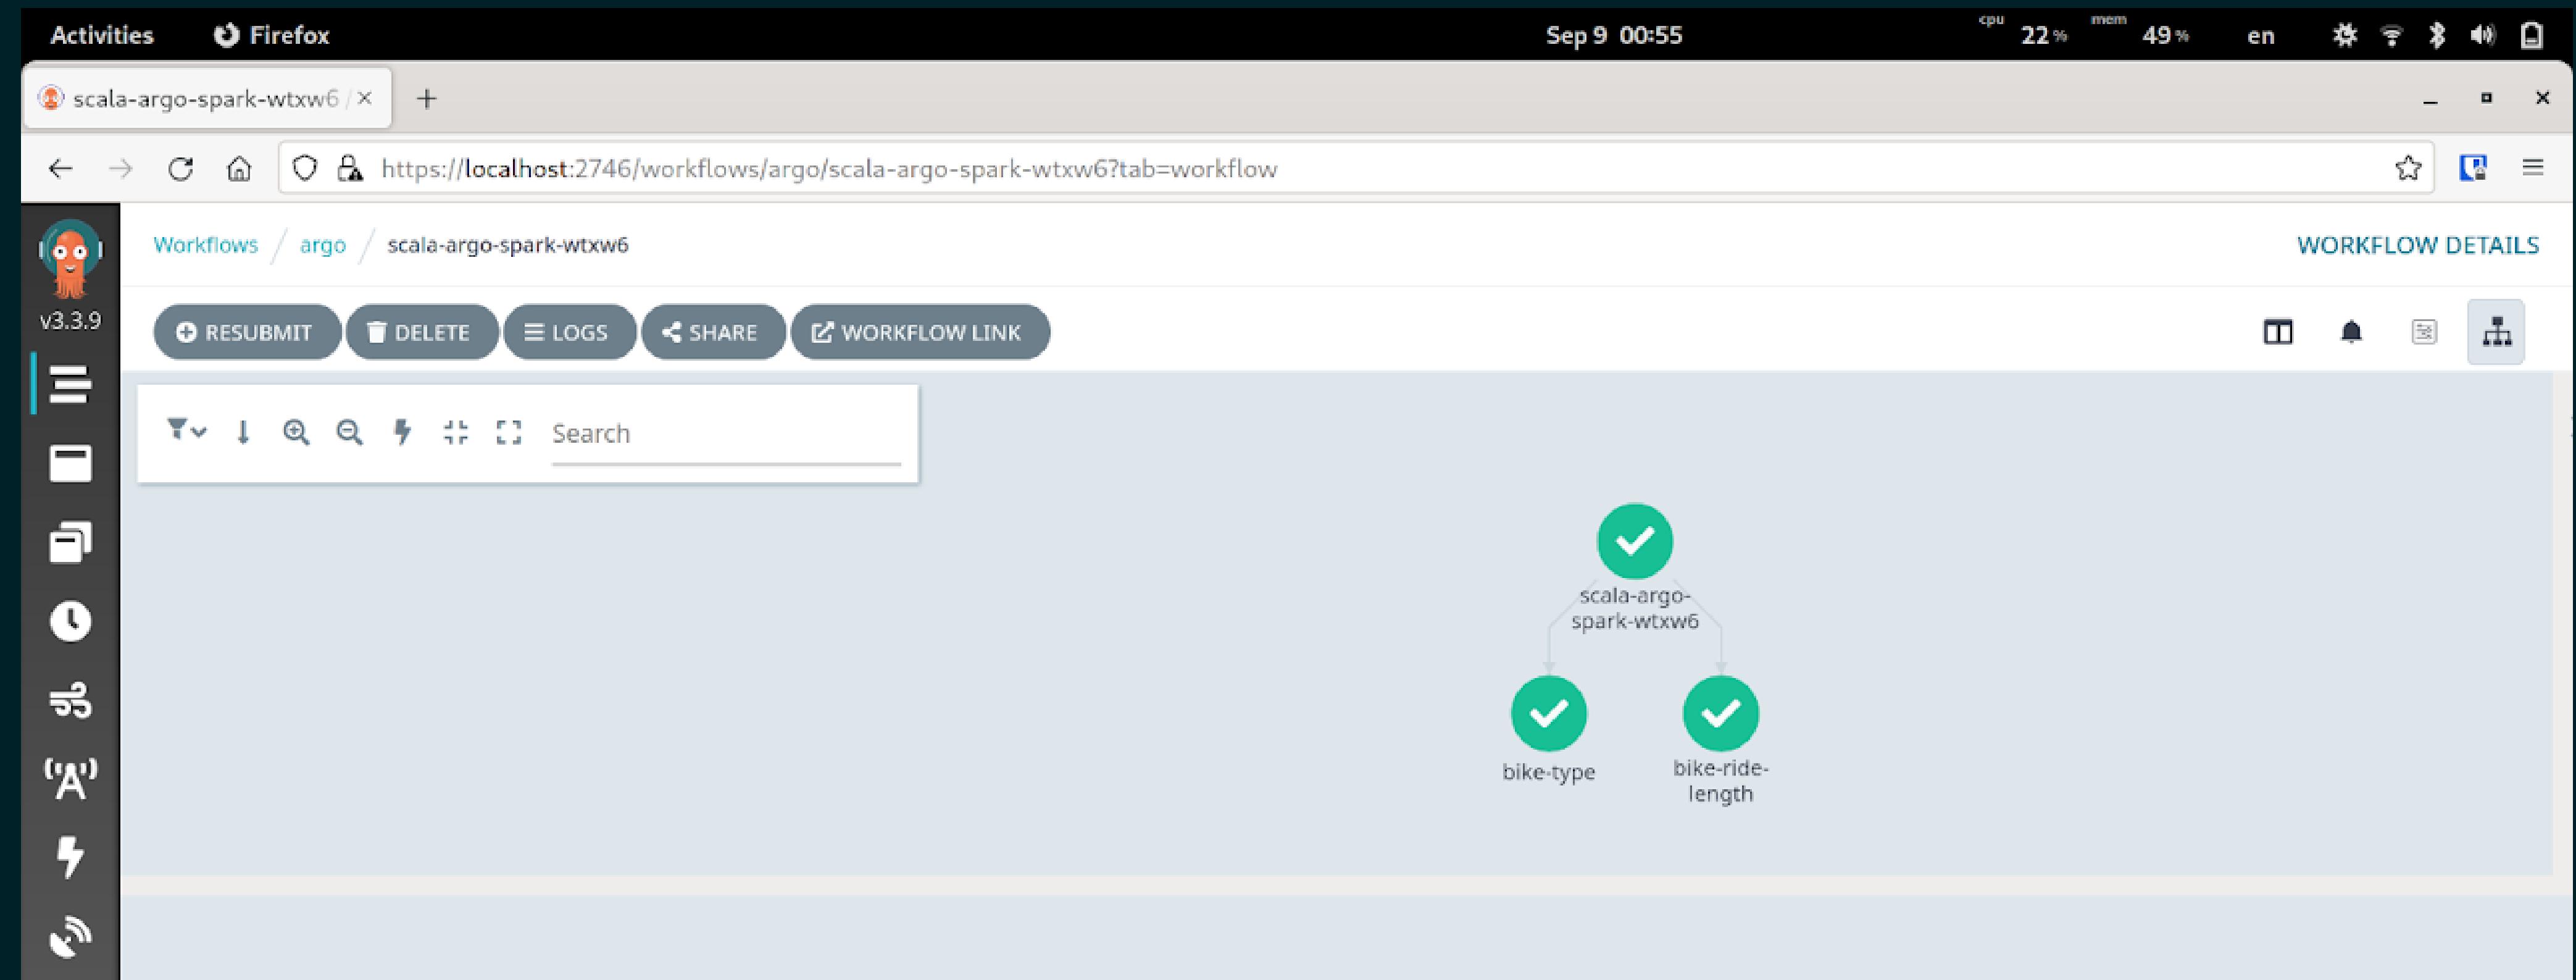
```
resource: |
  action: create
  successCondition: status.applicationState.state == COMPLETED
  failureCondition: status.applicationState.state == FAILED
manifest: |
  apiVersion: "sparkoperator.k8s.io/v1beta2"
  kind: SparkApplication
  metadata:
    generateName: spark-bike-type-
    namespace: default
  spec:
    type: Scala
    mode: cluster
    image: "argo-spark-integration-scala:latest"
    imagePullPolicy: Never
    mainClass: io.pipekit.oss.spark.argo.BikeTypeJob
    mainApplicationFile: "local:///app/bike.jar"
    sparkVersion: "3.1.1"
    restartPolicy:
      type: Never
    driver:
      coreRequest: "200m"
      coreLimit: "99999"
      memory: "1024m"
      serviceAccount: spark-driver
      labels:
        version: 3.1.1
    executor:
      coreRequest: "200m"
      coreLimit: "99999"
      instances: 2
      memory: "1024m"
      labels:
```



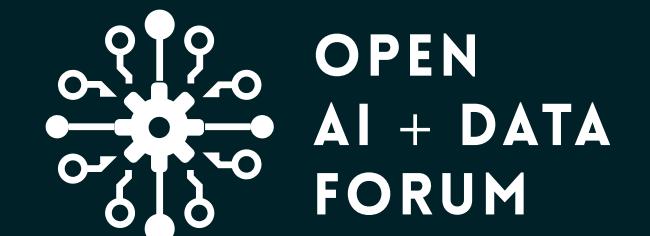
#OSSUMMIT

DEMO & HOW-TO

A successfully ran Spark job in the Argo UI:



Next steps

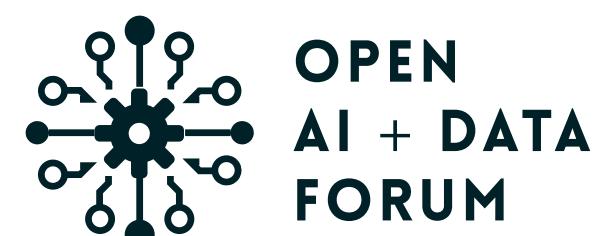


#OSSUMMIT

NEXT STEPS

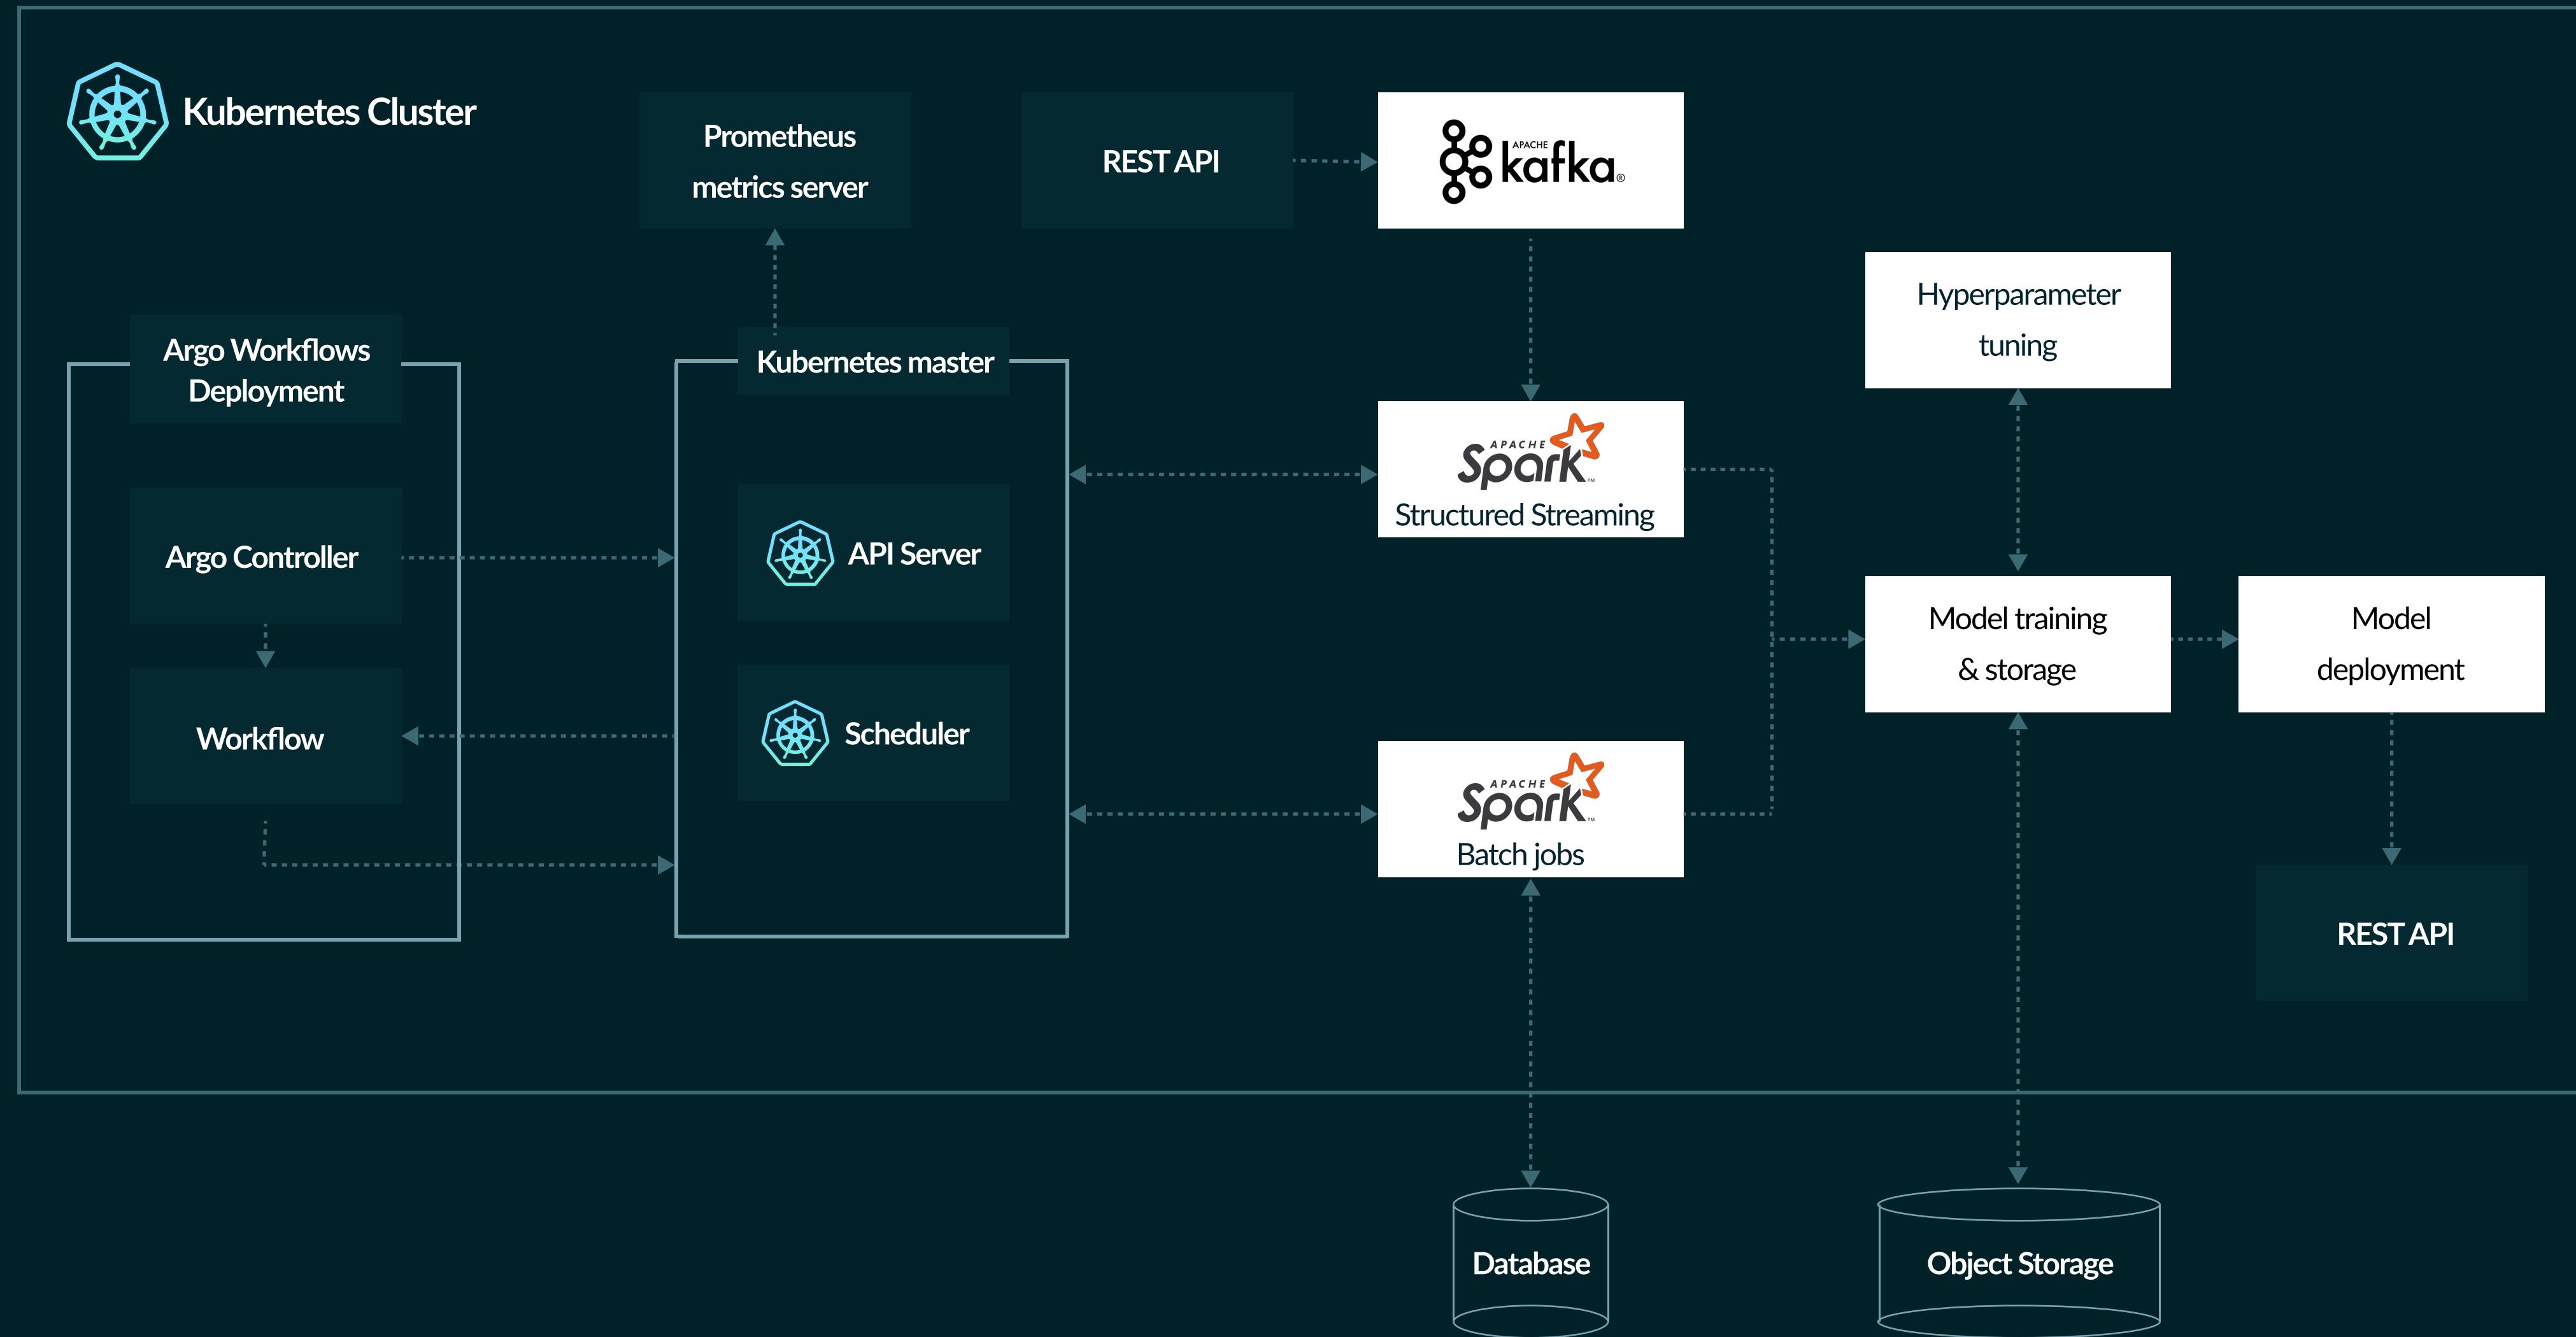
→ Now that we are running Spark on K8s,
what can we do next?

- Enable CRON workflows
- Define different node pools
- Cluster autoscaling - dynamically scale your Spark cluster on K8s
- Extend your data or ML pipeline with more capabilities

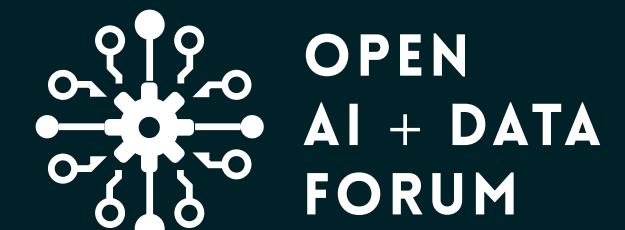


#OSSUMMIT

NEXT STEPS



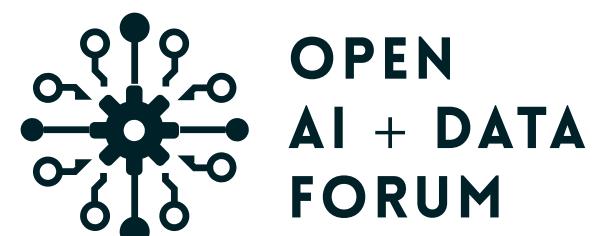
Resources & References



#OSSUMMIT

RESOURCES & REFERENCES

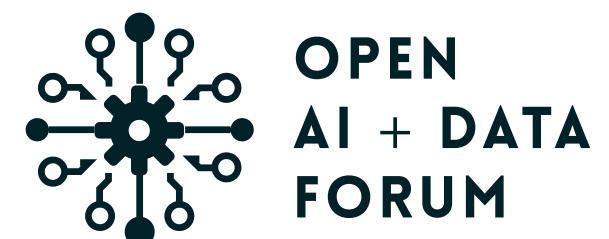
- ➡ Visit our repo: <https://github.com/pipekit/talk-demos/tree/main/oss-2022>
- ➡ Visit our blog: <https://pipekit.io/blog/>
- ➡ Follow [@pipekitio](#) on Twitter
- ➡ [Spark on K8s Operator GitHub repo](#)
- ➡ [Kubernetes resource management with Argo Workflows docs](#)



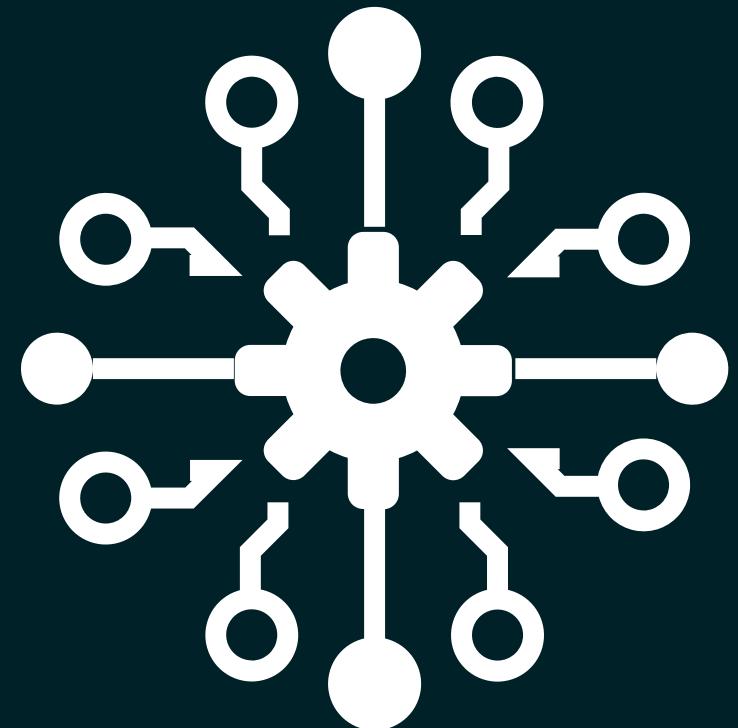
APPENDIX: ARGO WORKFLOWS VS. APACHE AIRFLOW

Argo workflows compared to Apache Airflow

- Lighter weight
- Container-native
- Dynamic, programmatic data pipelines out-of-the-box
- What about Python? Argo has a Python SDK, called [Hera](#)



#OSSUMMIT



OPEN AI + DATA FORUM



THE LINUX FOUNDATION

